

# Introduction to PyTorch

**Yifei Liu**

Tutorial session, STAT8056 2023 Spring

Slides and notebook available at: <https://github.com/yifei-liu-stat/stat8056-intro-pytorch>

For questions/discussions/typos, contact: liu00980@umn.edu



UNIVERSITY OF MINNESOTA

**Driven to Discover®**

# Outline

- Set up working environment
  - Coding platform
  - Practice: create conda environment for your project
- Introduction to PyTorch framework
- Deep learning with PyTorch
- Additional resources



# Coding platform

## On your local machine

### Use Anaconda to manage packages:

- Go to [anaconda.com](https://anaconda.com)
- Download Anaconda
- Open Anaconda Navigator (GUI)
  - Create an isolated environment for each of your projects.
  - Manage dependencies/versions of packages within the environment.
- Alternatively, one can use Anaconda in command-line style.

### Notebook-style coding:

- In Anaconda Navigator, install Jupyter Notebook (if not already)
- Launch Jupyter Notebook.

### Script-style coding:

- Choose an IDE to start with.
- Take VS Code for example:
  - Download [VS Code](https://code.visualstudio.com/).
  - Follow this [tutorial](#) to learn how to use Python in Visual Studio Code.



# Coding platform

## On a remote host/server

### Google Colab (for this tutorial)

- Notebook-style cloud computing.
  - <https://colab.research.google.com/>
  - Free version runs up to 12 hrs.
- Free GPU resource:
  - Runtime | Change runtime type | Hardware accelerator | GPU
- Many libraries (numpy, torch, scikit-learn ...) are pre-installed.
- Co-edit/code with collaborators.
- In sync. with your Google Drive.

## Resources from MSI:

- Need to be in a research group to use the cluster. [MSI access](#).
- Support Jupyter notebook now:
  - Connect to [UMN VPN](#) (off-campus only)
  - Visit [notebooks.msi.umn.edu](https://notebooks.msi.umn.edu)
  - Choose and start a server.
- [Free MSI tutorial \(with recordings\)](#)

## Other cloud services:

- [Deep-learning-in-cloud](#)
  - A list of cloud vendors either with free or paid services (with some free credits)



# UMN computing resources

MSI: <https://www.msi.umn.edu/>

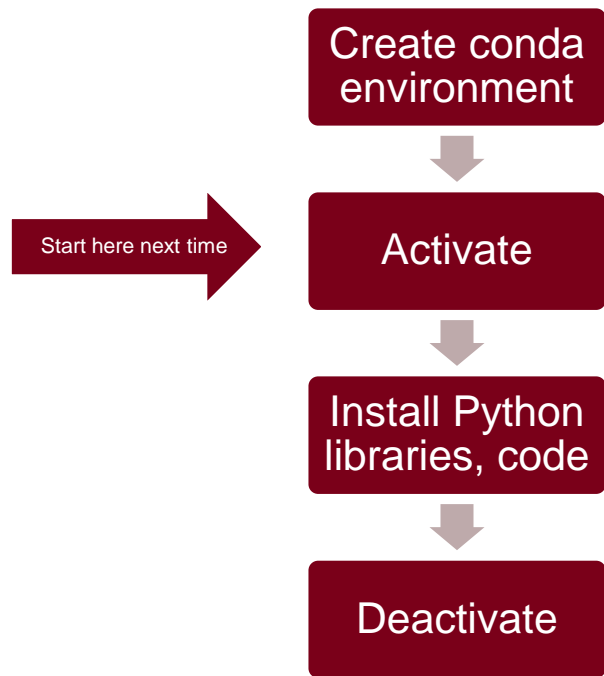
<a href="#">Introduction to Minnesota Supercomputing Institute (MSI)</a>	This tutorial is geared to new MSI users and will provide a highlevel introduction to the facilities and computational resources at MSI.	02/01/2022
<a href="#">Introduction to Linux</a>	This tutorial will provide an introduction to the Linux operating system, with particular attention paid to working from the command line	02/03/2022
<a href="#">Programming with Python</a>	Introduction to fundamentals of programming using the python language.	02/08/2022
<a href="#">Job Submission and Scheduling at MSI</a>	This tutorial will introduce users to MSI supercomputers, and provide an overview of how to submit calculations to the job schedulers	02/10/2022
<a href="#">Interactive Computing at MSI</a>	This two part tutorial will introduce you to the concept of interactive high performance computing, and provide attendees hands-on experience running interactive parallel jobs on the Mesabi HPC	02/15/2022
<a href="#">Data Storage Systems and Data Analysis Workflows for Research</a>	In this tutorial you will learn about the data storage systems available for academic research at the University of Minnesota	02/17/2022
<a href="#">Compiling and Debugging at MSI</a>	This tutorial will help users learn the basics of compiling and debugging their code on MSI systems	02/24/2022
<a href="#">Python for Scientific Computing</a>	This session includes efficient data processing with NumPy and Scipy, data visualization, and techniques for using python to drive parallel supercomputing tasks.	03/01/2022
<a href="#">RNA-Seq Analysis</a>	This tutorial covers the basics of differential expression analysis and touches on other RNA-seq topics such as transcriptome assembly.	03/03/2022
<a href="#">Parallel Computing On Agate</a>	This tutorial will help users learn the basics of parallel computation methods, including strategies for collecting calculations together for parallel execution.	03/31/2022

LATIS: <https://latisresearch.umn.edu/>

Feb. 18th   10:00am-noon	<a href="#">Creating Publication Worthy Visualizations without Code</a>	<a href="#">Registration</a>
Feb. 25th   10:00am-noon	<a href="#">Introduction to Computational Text Analysis</a>	<a href="#">Registration</a>
Mar. 4th   10:00am-noon	<a href="#">Reproducible research practices in Excel (yes, Excel)</a>	<a href="#">Registration</a>
Mar. 9th   10:00am-noon	<a href="#">Data Management in transition: Strategies for when you graduate</a>	<a href="#">Registration</a>
Mar. 18th   10:00am-noon <b>[RESCHEDULED]</b>	<a href="#">Advanced Nvivo</a>	<a href="#">Registration</a>
Mar. 25th   10:00am-noon	<a href="#">Introduction to parallel computing</a>	<a href="#">Registration</a>
April 1st   10:00am-noon	<a href="#">Introduction to SQL and Research Databases</a>	<a href="#">Registration</a>



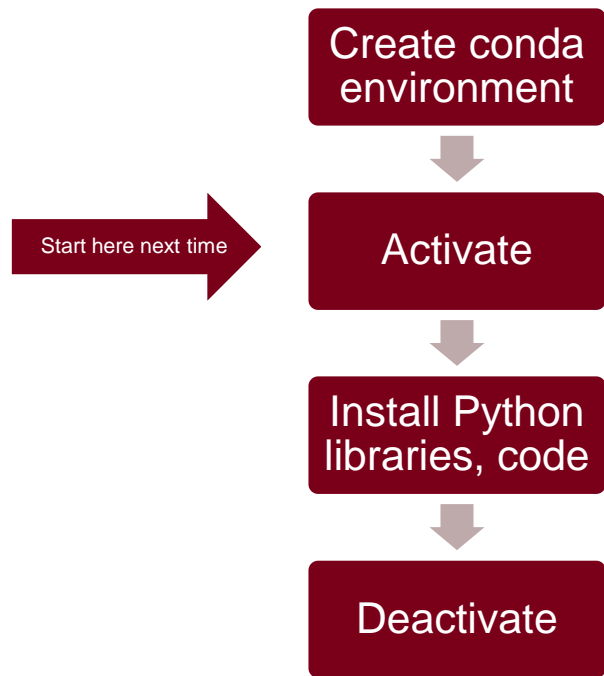
## Create a conda environment for your Python project



- An alternative way to Virtualenv
- Manage dependencies (**different versions** of Python and some libraries) **within an isolated project**
  - An example: Python library [foolbox](#)
- Key commands:
  - `conda create -n <envname> python=3.8.3`
  - `conda activate <envname>`
  - `conda install <package>=<version>`
  - `conda deactivate`
- Check [Conda Cheat Sheet](#)



## A hands-on exercise



- Make sure Anaconda is installed (either on your local machine or a remote host)
  - Windows: open Anaconda Prompt
  - MacOS/Linux: open terminal
  - Type `conda --version` to check availability
- Create an environment for course project
  - `conda create -n 8056proj python=3.9.13`
  - `conda activate 8056proj`
  - `conda install numpy`
  - `conda install jupyter notebook`
  - Do some coding:
    - Launch Jupyter Notebook: `jupyter notebook`
    - Launch Python terminal: `python`
  - `conda deactivate`



# Installation of PyTorch

- Choose configurations and install PyTorch from <http://pytorch.org/>

PyTorch Build	Stable (1.11.0)	Preview (Nightly)	LTS (1.8.2)	
Your OS	Linux	Mac	Windows	
Package	Conda	Pip	LibTorch	Source
Language	Python	C++ / Java		
Compute Platform	CUDA 10.2	CUDA 11.3	ROCm 4.2 (beta)	CPU
Run this Command:	<code>conda install pytorch torchvision torchaudio cudatoolkit=11.3 -c pytorch</code>			

Previous versions of PyTorch >

Run this within your conda environment

- If you use Google Colab, PyTorch is pre-installed with suitable configurations.





# Outline

- Set up working environment
- Introduction to PyTorch framework
  - Tensor, gradient and computation graph
  - Tensor manipulation on GPU
  - Use PyTorch as a general ML framework
- Deep learning with PyTorch
- Additional resources



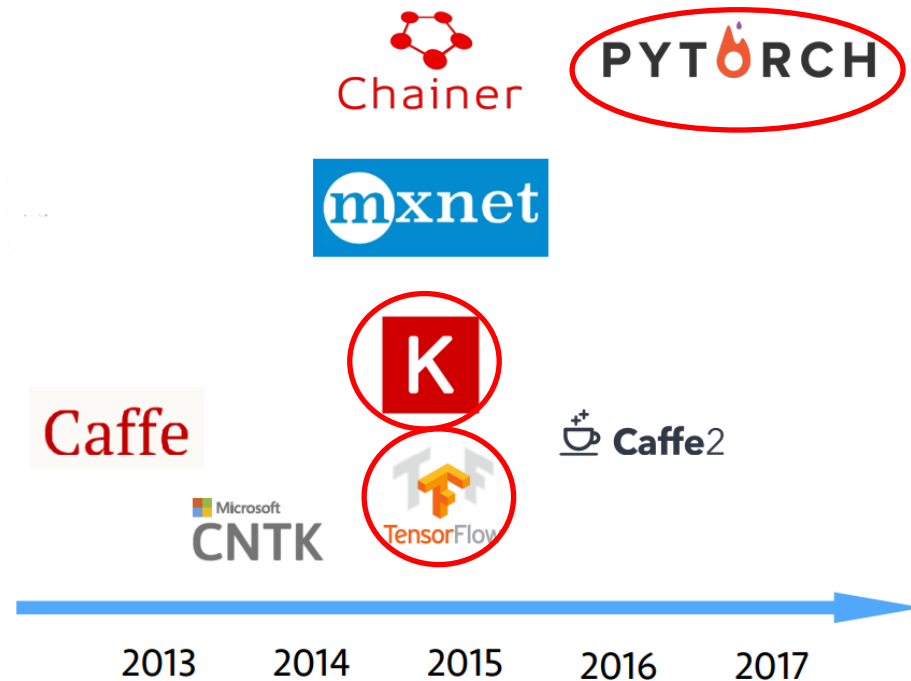
# Overview of PyTorch

A fancy version of NumPy but can:

- Trace gradient via autograd
- Accelerate via GPU
- Accommodate a bunch of machine learning functionalities including but not limited to neural network

More importantly,

- Easy interface compared to others
- Easy to understand and debug
- Check this [Google Trends](#)



Credit: "Gluon: new MXNet interface to accelerate research"



# Tensor

## What is a tensor?

- A multidimensional array
- Some examples:
  - 0D-tensor: scalar
  - 1D-tensor: vector
  - 2D-tensor: matrix; non-RGB image
  - 3D-tensor: RGB image (3 channels)
  - 4D-tensor: **one** RGB video clip
  - 5D-tensor: **a collection of** RGB video clips
  - 6D-tensor: ???
- Tensor in PyTorch:
  - Just like arrays in Numpy
  - `t.numpy()` and `torch.from_numpy(a)`

## np.array versus torch.tensor

- Similar object creations:
  - All-ones, all-zeros, identity matrix, random...
  - Check [PyTorch tensor creations](#)
- Similar math operations:
  - Indexing, slicing, reshape, transpose, tensor product, element-wise operation...
  - Check [PyTorch tensor math operations](#)

```
import numpy as np
myarray = np.ones(3)
print(myarray)
print(myarray + 1.5)
```

```
[1.  1.  1.]
[2.5  2.5  2.5]
```

```
import torch
mytensor = torch.ones(3)
print(mytensor)
print(mytensor + 1.5)
```

```
tensor([1., 1., 1.])
tensor([2.5000, 2.5000, 2.5000])
```



# Tensor (cont.)

- Enable gradient tracing:
  - Upon any tensor creation, set `requires_grad = True`
  - For any existent tensor `t`, call `t.requires_grad_(True)`
- Calculate gradient:
  - Do some calculation from `t` and get `f`
  - Call `f.backward()` for calculation
  - Call `t.grad` to access the gradient
- Disable gradient tracing:
  - (Permanently) `t = t.detach()`
  - (Temporarily) with `torch.no_grad()`:

enable  
gradient

do some  
calculation

call  
.backward()

access the  
gradient



```
# create tensors, and enable gradient tracing
c = torch.tensor(1.)
t = torch.tensor(3., requires_grad = True)
print("Tensor c:", c)
print("Tensor t:", t)

# do some calculation
s = t**2
f = 3 * s + c

# calculate the gradient of f w.r.t. t
f.backward()

# access the gradient
print("Gradient df/dt:", t.grad)
```



```
Tensor c: tensor(1.)
Tensor t: tensor(3., requires_grad=True)
Gradient df/dt: tensor(18.)
```



# Dynamic computation graph

t

c

enable  
gradient

do some  
calculation

call  
.backward()

access the  
gradient



```
# create tensors, and enable gradient tracing
c = torch.tensor(1.)
t = torch.tensor(3., requires_grad = True)
print("Tensor c:", c)
print("Tensor t:", t)

# do some calculation
s = t**2
f = 3 * s + c

# calculate the gradient of f w.r.t. t
f.backward()

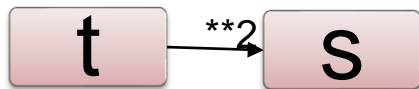
# access the gradient
print("Gradient df/dt:", t.grad)
```



```
Tensor c: tensor(1.)
Tensor t: tensor(3., requires_grad=True)
Gradient df/dt: tensor(18.)
```



# Dynamic computation graph



```
# create tensors, and enable gradient tracing
c = torch.tensor(1.)
t = torch.tensor(3., requires_grad = True)
print("Tensor c:", c)
print("Tensor t:", t)

# do some calculation
s = t**2
f = 3 * s + c

# calculate the gradient of f w.r.t. t
f.backward()

# access the gradient
print("Gradient df/dt:", t.grad)
```

Tensor c: tensor(1.)  
Tensor t: tensor(3., requires\_grad=True)  
Gradient df/dt: tensor(18.)



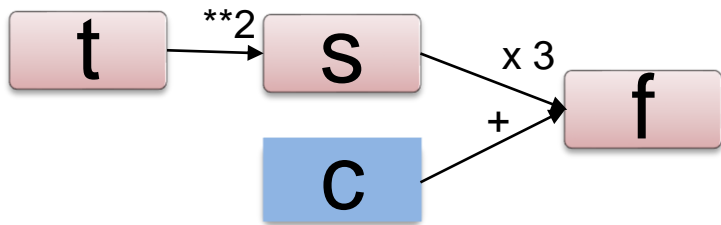
# Dynamic computation graph

enable  
gradient

do some  
calculation

call  
.backward()

access the  
gradient



```
# create tensors, and enable gradient tracing
c = torch.tensor(1.)
t = torch.tensor(3., requires_grad = True)
print("Tensor c:", c)
print("Tensor t:", t)

# do some calculation
s = t**2
f = 3 * s + c

# calculate the gradient of f w.r.t. t
f.backward()

# access the gradient
print("Gradient df/dt:", t.grad)
```



```
Tensor c: tensor(1.)
Tensor t: tensor(3., requires_grad=True)
Gradient df/dt: tensor(18.)
```



# Dynamic computation graph

enable  
gradient

do some  
calculation

call  
.backward()

access the  
gradient

The graph structure is destroyed  
Once you call `.backward()`

t

s

c

f



```
# create tensors, and enable gradient tracing
c = torch.tensor(1.)
t = torch.tensor(3., requires_grad = True)
print("Tensor c:", c)
print("Tensor t:", t)

# do some calculation
s = t**2
f = 3 * s + c

# calculate the gradient of f w.r.t. t
f.backward()

# access the gradient
print("Gradient df/dt:", t.grad)
```

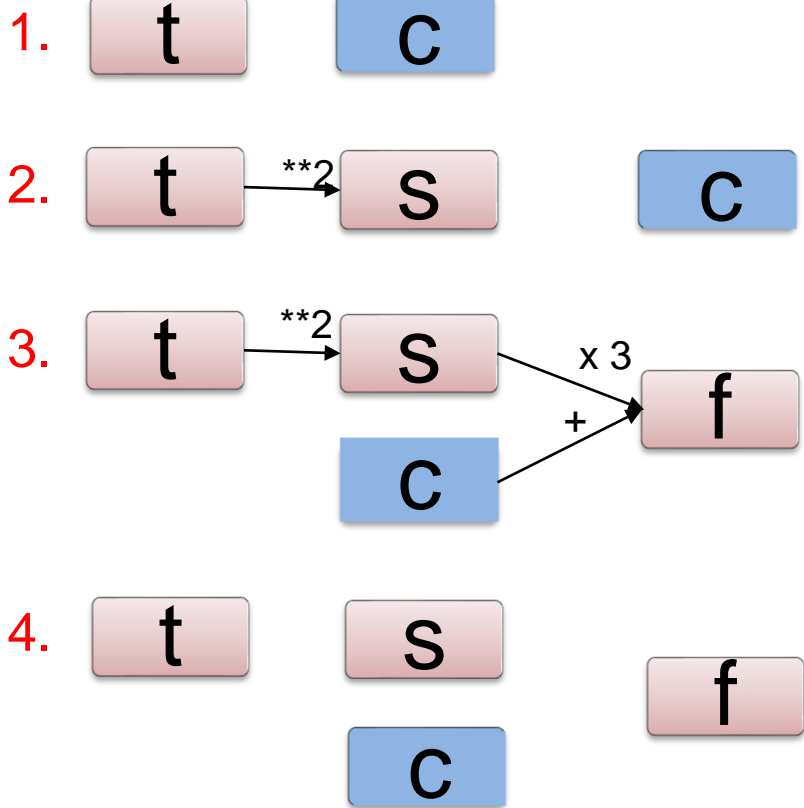


```
Tensor c: tensor(1.)
Tensor t: tensor(3., requires_grad=True)
Gradient df/dt: tensor(18.)
```





# Dynamic computation graph



```
# create tensors, and enable gradient tracing
c = torch.tensor(1.)
t = torch.tensor(3., requires_grad = True)
print("Tensor c:", c)
print("Tensor t:", t)

# do some calculation
s = t**2
f = 3 * s + c

# calculate the gradient of f w.r.t. t
f.backward()

# access the gradient
print("Gradient df/dt:", t.grad)
```



```
Tensor c: tensor(1.)
Tensor t: tensor(3., requires_grad=True)
Gradient df/dt: tensor(18.)
```



# Why should we understand PyTorch computation graph?

- Essentially, PyTorch is using **chain rule** to calculate the gradient
- A computation graph defines how the chain rule applies to your calculation
- **Common MISTAKES:**
  - Call `.backward()` when there is no graph
  - Retrieve gradient of non-leaf nodes
  - Gradient accumulation. To solve this problem, call `t.grad.zero_()` before building the second graph

## Illustration of gradient accumulation

```
# first back propagation
t = torch.tensor(3., requires_grad = True)
s = t**2
s.backward()
print("Gradient ds/dt:", t.grad)

# second back propagation
f = 5 * t + 1
f.backward()
print("Gradient df/dt (without emptying t.grad):", t.grad)
```

Gradient ds/dt: tensor(6.)  
Gradient df/dt (without emptying t.grad): **tensor(11.)** should be 5

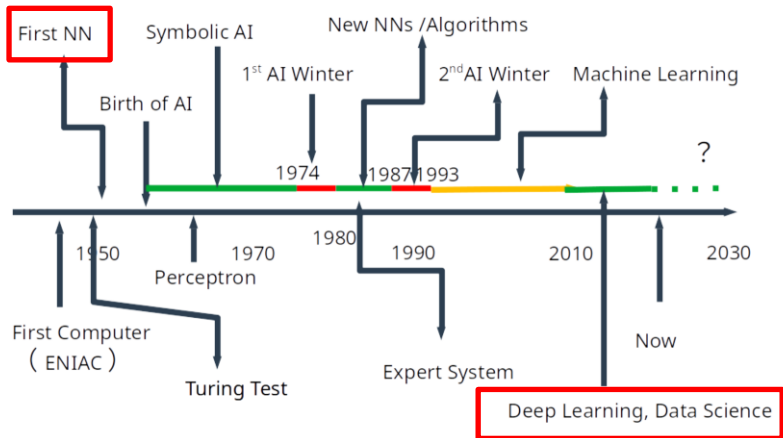
```
t.grad.zero_() # zero out the gradient
f = 5 * t + 1
f.backward()
print("Gradient df/dt (after emptying t.grad):", t.grad)
```

Gradient df/dt (after emptying t.grad): **tensor(5.)**



# Calculation on GPU

- Neural network dates back to 1950s, but became popular **only recently (last decade)** due to huge improvements in computation



Credit: CSCI 5980/8980: Think Deep Learning offered by Ju Sun, UMN

- GPU allows faster **large-scale matrix multiplication** (or tensor product)
- Use GPU on PyTorch:
  - `torch.cuda.is_available()` # GPU?
  - `t = t.cuda()` # move to GPU
  - `t = t.cpu()` # move to CPU
  - `t.device` # cpu or cuda: index
  - The transfer time can be long
  - Make sure all tensor manipulations are performed on the same device
- Advanced: [train on multiple GPUs](#)



# Time comparison: matrix multiplication on CPU and GPU

- We want to perform matrix multiplication  $C = AB$  where both A and B are 3000 x 3000 matrices
- Enable GPU on [Google Colab](#) (free!)
  1. Runtime
  2. Change runtime type
  3. Hardware Accelerator: GPU
  4. Runtime -> Restart Runtime
- Much faster on GPU for large problem!

```
np.random.seed(8056)

d = 3000

# C = A B with numpy on CPU
A = np.random.rand(d, d)
B = np.random.rand(d, d)
begin = time.time()
C = A.dot(B)
print(f"CPU time (numpy): {time.time() - begin} s")

# C = A B with torch on CPU
A_t_cpu = torch.tensor(A)
B_t_cpu = torch.tensor(B)
begin = time.time()
C = torch.mm(A_t_cpu, B_t_cpu)
print(f"CPU time (torch): {time.time() - begin} s")

# C = A B with torch on GPU
# make sure GPU is available
A_t_gpu = torch.tensor(A).cuda()
B_t_gpu = torch.tensor(B).cuda()
begin = time.time()
C = torch.mm(A_t_gpu, B_t_gpu)
print(f"GPU time (torch): {time.time() - begin} s")
```

CPU time (numpy): 1.0066237449645996 s  
CPU time (torch): 0.8800492286682129 s  
GPU time (torch): 0.0038247108459472656 s



# GPU not always better

- Calculation on GPU should be optimized smartly since
  - GPU has **limited memory** compared to CPU
  - **Transfer/overhead time** from CPU to GPU can be large

```
▶ d = 3000
begin = time.time()
A = torch.rand((d, d)).cuda()
B = torch.rand((d, d)).cuda()
print(f"Time from CPU to GPU: {time.time() - begin} s")
begin = time.time()
C = torch.mm(A, B)
print(f"GPU time (torch): {time.time() - begin} s")

☞ Time from CPU to GPU: 0.15282464027404785 s
GPU time (torch): 0.0008263587951660156 s
```

NVIDIA-SMI 440.100			Driver Version: 440.100			CUDA Version: 10.2		
GPU	Name	Persistence-M	Bus-Id	Disp.A	Volatile Uncorr. ECC			
Fan	Temp	Perf	Pwr:Usage/Cap	Memory-Usage	GPU-Util	Compute	M.	
0	TITAN RTX	Off	00000000:1A:00.0	Off	N/A			
41%	56C	P2	129W / 280W	13556MiB / 24220MiB	61%	Default		
1	TITAN RTX	Off	00000000:1B:00.0	Off	N/A			
41%	55C	P2	127W / 280W	18798MiB / 24220MiB	60%	Default		
2	TITAN RTX	Off	00000000:3D:00.0	Off	N/A			
41%	48C	P2	121W / 280W	13616MiB / 24220MiB	62%	Default		
3	TITAN RTX	Off	00000000:3E:00.0	Off	N/A			
41%	56C	P2	121W / 280W	20962MiB / 24220MiB	61%	Default		
4	TITAN RTX	Off	00000000:88:00.0	Off	N/A			
41%	53C	P2	162W / 280W	13590MiB / 24220MiB	92%	Default		
5	TITAN RTX	Off	00000000:89:00.0	Off	N/A			
41%	59C	P2	155W / 280W	18832MiB / 24220MiB	61%	Default		
6	TITAN RTX	Off	00000000:B1:00.0	Off	N/A			
41%	55C	P2	169W / 280W	13650MiB / 24220MiB	90%	Default		
7	TITAN RTX	Off	00000000:B2:00.0	Off	N/A			
41%	56C	P2	109W / 280W	20996MiB / 24220MiB	91%	Default		



# ML in PyTorch

- PyTorch provides a framework for general ML modeling **including but not limited to deep learning**
- What is a common ML pipeline?
  - Prepare your dataset
  - Choose a training model
  - Choose a loss/objective function
  - **Optimization** (calculation; gradient)
  - **Evaluation** and inference
- The go-to optimization: SGD
  - Instead of using gradient calculated from all training samples, we **only use the gradient from a randomly chosen sample**
- Practical one: (Mini-) Batch SGD
  - Use gradient from **a batch of samples**
  - **Batch size**: # samples in one batch
  - **Epoch**: **a full pass** of all training samples
  - Special cases:
    - Batch size = 1: vanilla SGD
    - Batch size = n: GD



# Linear regression with SGD

## – The naive way

- The true model:  $y = 2 * x + 1$
- Some other setups:
  - $n = 2000$ ,  $x \sim \text{Uniform}(0, 1)$
  - Number of epochs: 10
  - Batch size: 200 (10 batches per epoch)
  - Learning rate: 0.05
  - Initialize both intercept and slope with  $\text{Uniform}(0, 1)$



```
n = 2000
x = torch.rand(n)

# create dataset with true model
a0 = 2
b0 = 1
y = a0 * x + b0

# set up optimization parameter of SGD
a = torch.rand(1, requires_grad = True)
b = torch.rand(1, requires_grad = True)
nepochs = 10
batch_size = 200
lr = 0.5
```



```

▶ for epoch in range(nepochs):
    for batch in range(round(n / batch_size)):
        start = batch * batch_size
        end = start + batch_size
        # perform update on a batch
        x_batch = x[start:end]
        y_batch = y[start:end]

        # build computation graph
        y_hat = a * x_batch + b
        myloss = torch.mean((y_batch - y_hat)**2)

        # gradient calculation
        myloss.backward()

        # SGD update
        with torch.no_grad():
            a -= lr * a.grad
            b -= lr * b.grad

        # avoid gradient accumulation
        a.grad.zero_()
        b.grad.zero_()

    print(f"Epoch: {epoch + 1} / {nepochs}")
    print(f"MSE: {myloss: .2e}; a: {a.item():.4f}; b: {b.item():.4f}")

```

Outer iteration of SGD (epoch)

## Linear regression with SGD – The naive way

### Output:

```

❏ Epoch: 1 / 10
MSE: 3.69e-02; a: 1.4460; b: 1.2822
Epoch: 2 / 10
MSE: 9.24e-03; a: 1.7226; b: 1.1413
Epoch: 3 / 10
MSE: 2.32e-03; a: 1.8611; b: 1.0707
Epoch: 4 / 10
MSE: 5.81e-04; a: 1.9305; b: 1.0354
Epoch: 5 / 10
MSE: 1.46e-04; a: 1.9652; b: 1.0177
Epoch: 6 / 10
MSE: 3.65e-05; a: 1.9826; b: 1.0089
Epoch: 7 / 10
MSE: 9.15e-06; a: 1.9913; b: 1.0044
Epoch: 8 / 10
MSE: 2.29e-06; a: 1.9956; b: 1.0022
Epoch: 9 / 10
MSE: 5.75e-07; a: 1.9978; b: 1.0011
Epoch: 10 / 10
MSE: 1.44e-07; a: 1.9989; b: 1.0006

```







```
for epoch in range(nepochs):  
    for batch in range(round(n / batch_size)):  
        start = batch * batch_size  
        end = start + batch_size  
        # perform update on a batch  
        x_batch = x[start:end]  
        y_batch = y[start:end]  
  
        # build computation graph  
        y_hat = a * x_batch + b  
        myloss = torch.mean((y_batch - y_hat)**2)  
  
        # gradient calculation  
        myloss.backward()  
  
        # SGD update  
        with torch.no_grad():  
            a -= lr * a.grad  
            b -= lr * b.grad  
  
        # avoid gradient accumulation  
        a.grad.zero_()  
        b.grad.zero_()  
  
print(f"Epoch: {epoch + 1} / {nepochs}")  
print(f"MSE: {myloss: .2e}; a: {a.item():.4f}; b: {b.item():.4f}")
```

Inner iteration  
of SGD (batch)

## Linear regression with SGD – The naive way

### Output:

```
Epoch: 1 / 10  
MSE: 3.69e-02; a: 1.4460; b: 1.2822  
Epoch: 2 / 10  
MSE: 9.24e-03; a: 1.7226; b: 1.1413  
Epoch: 3 / 10  
MSE: 2.32e-03; a: 1.8611; b: 1.0707  
Epoch: 4 / 10  
MSE: 5.81e-04; a: 1.9305; b: 1.0354  
Epoch: 5 / 10  
MSE: 1.46e-04; a: 1.9652; b: 1.0177  
Epoch: 6 / 10  
MSE: 3.65e-05; a: 1.9826; b: 1.0089  
Epoch: 7 / 10  
MSE: 9.15e-06; a: 1.9913; b: 1.0044  
Epoch: 8 / 10  
MSE: 2.29e-06; a: 1.9956; b: 1.0022  
Epoch: 9 / 10  
MSE: 5.75e-07; a: 1.9978; b: 1.0011  
Epoch: 10 / 10  
MSE: 1.44e-07; a: 1.9989; b: 1.0006
```





```
for epoch in range(nepochs):
    for batch in range(round(n / batch_size)):
        start = batch * batch_size
        end = start + batch_size
        # perform update on a batch
        x_batch = x[start:end]
        y_batch = y[start:end]

        # build computation graph
        y_hat = a * x_batch + b
        myloss = torch.mean((y_batch - y_hat)**2)

        # gradient calculation
        myloss.backward()

        # SGD update
        with torch.no_grad():
            a -= lr * a.grad
            b -= lr * b.grad

        # avoid gradient accumulation
        a.grad.zero_()
        b.grad.zero_()

print(f"Epoch: {epoch + 1} / {nepochs}")
print(f"MSE: {myloss: .2e}; a: {a.item():.4f}; b: {b.item():.4f}")
```

1. Use a batch of data to  
build a computation graph

## Linear regression with SGD – The naive way

### Output:

```
Epoch: 1 / 10
MSE: 3.69e-02; a: 1.4460; b: 1.2822
Epoch: 2 / 10
MSE: 9.24e-03; a: 1.7226; b: 1.1413
Epoch: 3 / 10
MSE: 2.32e-03; a: 1.8611; b: 1.0707
Epoch: 4 / 10
MSE: 5.81e-04; a: 1.9305; b: 1.0354
Epoch: 5 / 10
MSE: 1.46e-04; a: 1.9652; b: 1.0177
Epoch: 6 / 10
MSE: 3.65e-05; a: 1.9826; b: 1.0089
Epoch: 7 / 10
MSE: 9.15e-06; a: 1.9913; b: 1.0044
Epoch: 8 / 10
MSE: 2.29e-06; a: 1.9956; b: 1.0022
Epoch: 9 / 10
MSE: 5.75e-07; a: 1.9978; b: 1.0011
Epoch: 10 / 10
MSE: 1.44e-07; a: 1.9989; b: 1.0006
```





```
for epoch in range(nepochs):
    for batch in range(round(n / batch_size)):
        start = batch * batch_size
        end = start + batch_size
        # perform update on a batch
        x_batch = x[start:end]
        y_batch = y[start:end]

        # build computation graph
        y_hat = a * x_batch + b
        myloss = torch.mean((y_batch - y_hat)**2)

        # gradient calculation
        myloss.backward()

        # SGD update
        with torch.no_grad():
            a -= lr * a.grad
            b -= lr * b.grad

        # avoid gradient accumulation
        a.grad.zero_()
        b.grad.zero_()

print(f"Epoch: {epoch + 1} / {nepochs}")
print(f"MSE: {myloss: .2e}; a: {a.item():.4f}; b: {b.item():.4f}")
```

2. Calculate gradient based  
on the computation graph

## Linear regression with SGD – The naive way

### Output:

```
Epoch: 1 / 10
MSE: 3.69e-02; a: 1.4460; b: 1.2822
Epoch: 2 / 10
MSE: 9.24e-03; a: 1.7226; b: 1.1413
Epoch: 3 / 10
MSE: 2.32e-03; a: 1.8611; b: 1.0707
Epoch: 4 / 10
MSE: 5.81e-04; a: 1.9305; b: 1.0354
Epoch: 5 / 10
MSE: 1.46e-04; a: 1.9652; b: 1.0177
Epoch: 6 / 10
MSE: 3.65e-05; a: 1.9826; b: 1.0089
Epoch: 7 / 10
MSE: 9.15e-06; a: 1.9913; b: 1.0044
Epoch: 8 / 10
MSE: 2.29e-06; a: 1.9956; b: 1.0022
Epoch: 9 / 10
MSE: 5.75e-07; a: 1.9978; b: 1.0011
Epoch: 10 / 10
MSE: 1.44e-07; a: 1.9989; b: 1.0006
```





```
for epoch in range(nepochs):
    for batch in range(round(n / batch_size)):
        start = batch * batch_size
        end = start + batch_size
        # perform update on a batch
        x_batch = x[start:end]
        y_batch = y[start:end]

        # build computation graph
        y_hat = a * x_batch + b
        myloss = torch.mean((y_batch - y_hat)**2)

        # gradient calculation
        myloss.backward()

        # SGD update
        with torch.no_grad():
            a -= lr * a.grad
            b -= lr * b.grad

        # avoid gradient accumulation
        a.grad.zero_()
        b.grad.zero_()

print(f"Epoch: {epoch + 1} / {nepochs}")
print(f"MSE: {myloss: .2e}; a: {a.item():.4f}; b: {b.item():.4f}")
```

3. Update parameters  
and zero out gradients

## Linear regression with SGD – The naive way

### Output:

```
Epoch: 1 / 10
MSE: 3.69e-02; a: 1.4460; b: 1.2822
Epoch: 2 / 10
MSE: 9.24e-03; a: 1.7226; b: 1.1413
Epoch: 3 / 10
MSE: 2.32e-03; a: 1.8611; b: 1.0707
Epoch: 4 / 10
MSE: 5.81e-04; a: 1.9305; b: 1.0354
Epoch: 5 / 10
MSE: 1.46e-04; a: 1.9652; b: 1.0177
Epoch: 6 / 10
MSE: 3.65e-05; a: 1.9826; b: 1.0089
Epoch: 7 / 10
MSE: 9.15e-06; a: 1.9913; b: 1.0044
Epoch: 8 / 10
MSE: 2.29e-06; a: 1.9956; b: 1.0022
Epoch: 9 / 10
MSE: 5.75e-07; a: 1.9978; b: 1.0011
Epoch: 10 / 10
MSE: 1.44e-07; a: 1.9989; b: 1.0006
```



# Linear regression with SGD

## – The PyTorch way

### Built-in functionalities

- Dataset preparation:
  - [torch.utils.data](#)
  - Check [Datasets & DataLoaders](#)
- Define training model:
  - Check [torch.nn](#) for all kinds of components to build your own model
- Optimization algorithms:
  - Check [torch.optim](#) for various opt methods

```
from torch import nn, optim
from torch.utils.data import Dataset, DataLoader
```

```
class SimpleLinear(Dataset):
    def __init__(self, n, a, b):
        self.n = n
        self.x = torch.rand(n, 1)
        self.y = a * self.x + b

    def __len__(self):
        return self.n

    def __getitem__(self, idx):
        return self.x[idx], self.y[idx]

# create a training model
# (cont.) also initializes a and b
mymodel = nn.Sequential(nn.Linear(1, 1))

# prepare dataset and dataloaders
mydata = SimpleLinear(n = 2000, a = 2, b = 1)
mydataloader = DataLoader(mydata, batch_size = 200)

# set up optimization with SGD
criterion = nn.MSELoss()
optimizer = optim.SGD(mymodel.parameters(), lr = 0.5)
```



# Linear regression with SGD

## – The PyTorch way

### Prepare our datasets

- Dataset (a Python class)
  - `__init__`: initialize the dataset
  - `__len__`: sample size of the dataset
  - `__getitem__`: fetch a sample with `idx`
  - As the input of `DataLoader` function
- DataLoader (a Python function)
  - Divide the dataset into batches
  - Check [this](#) for more advanced usages
  - Have some tricks to reduce overhead time when transferring data to GPU

```
class SimpleLinear(Dataset):
    def __init__(self, n, a, b):
        self.n = n
        self.x = torch.rand(n, 1)
        self.y = a * self.x + b

    def __len__(self):
        return self.n

    def __getitem__(self, idx):
        return self.x[idx], self.y[idx]

# create a training model
# (cont.) also initializes a and b
mymodel = nn.Sequential(nn.Linear(1, 1))

# prepare dataset and dataloaders
mydata = SimpleLinear(n = 2000, a = 2, b = 1)
mydataloader = DataLoader(mydata, batch_size = 200)

# set up optimization with SGD
criterion = nn.MSELoss()
optimizer = optim.SGD(mymodel.parameters(), lr = 0.5)
```



# Linear regression with SGD

## – The PyTorch way

Create the training model:

- `nn.Sequential`
  - Build a model with sequential operations
  - `nn.Linear(m, n)`
    - A **linear operator** of shape  $n \times m$
    - A **bias vector** of shape  $n \times 1$  (default)
  - All parameters are **initialized automatically upon creation**
  - Check [torch.nn](https://pytorch.org/docs/nn.html) for other operations



```
class SimpleLinear(Dataset):
    def __init__(self, n, a, b):
        self.n = n
        self.x = torch.rand(n, 1)
        self.y = a * self.x + b

    def __len__(self):
        return self.n

    def __getitem__(self, idx):
        return self.x[idx], self.y[idx]

# create a training model
# (cont.) also initializes a and b
mymodel = nn.Sequential(nn.Linear(1, 1))

# prepare dataset and dataloaders
mydata = SimpleLinear(n = 2000, a = 2, b = 1)
mydataloader = DataLoader(mydata, batch_size = 200)

# set up optimization with SGD
criterion = nn.MSELoss()
optimizer = optim.SGD(mymodel.parameters(), lr = 0.5)
```



# Linear regression with SGD

## – The PyTorch way

Set up optimization:

- The loss function
- The optimization method
- Optimization parameters

```
class SimpleLinear(Dataset):  
    def __init__(self, n, a, b):  
        self.n = n  
        self.x = torch.rand(n, 1)  
        self.y = a * self.x + b  
  
    def __len__(self):  
        return self.n  
  
    def __getitem__(self, idx):  
        return self.x[idx], self.y[idx]  
  
# create a training model  
# (cont.) also initializes a and b  
mymodel = nn.Sequential(nn.Linear(1, 1))  
  
# prepare dataset and dataloaders  
mydata = SimpleLinear(n = 2000, a = 2, b = 1)  
mydataloader = DataLoader(mydata, batch_size = 200)  
  
# set up optimization with SGD  
criterion = nn.MSELoss()  
optimizer = optim.SGD(mymodel.parameters(), lr = 0.5)
```





# Linear regression with SGD

## – The PyTorch way



```
nepochs = 10
for epoch in range(nepochs):
    for x_batch, y_batch in mydataloader:
        # build computation graph
        yhat = mymodel(x_batch)
        myloss = criterion(y_batch, yhat)

        # optimization
        optimizer.zero_grad() # zero out gradient
        myloss.backward() # back propagation
        optimizer.step() # update parameters
```

## Output:

```
Epoch: 1 / 10
MSE: 3.52e-02; a: 1.3904; b: 1.3368
Epoch: 2 / 10
MSE: 9.21e-03; a: 1.6882; b: 1.1722
Epoch: 3 / 10
MSE: 2.41e-03; a: 1.8406; b: 1.0881
Epoch: 4 / 10
MSE: 6.30e-04; a: 1.9185; b: 1.0450
Epoch: 5 / 10
MSE: 1.65e-04; a: 1.9583; b: 1.0230
Epoch: 6 / 10
MSE: 4.31e-05; a: 1.9787; b: 1.0118
Epoch: 7 / 10
MSE: 1.13e-05; a: 1.9891; b: 1.0060
Epoch: 8 / 10
MSE: 2.95e-06; a: 1.9944; b: 1.0031
Epoch: 9 / 10
MSE: 7.71e-07; a: 1.9971; b: 1.0016
Epoch: 10 / 10
MSE: 2.02e-07; a: 1.9985; b: 1.0008
```



# Linear regression with SGD

## – The PyTorch way



```
nepochs = 10
for epoch in range(nepochs):
    for x_batch, y_batch in mydataloader:
        # build computation graph
        yhat = mymodel(x_batch)
        myloss = criterion(y_batch, yhat)

        # optimization
        optimizer.zero_grad() # zero out gradient
        myloss.backward() # back propagation
        optimizer.step() # update parameters
```

The whole  
optimization part

## Output:

```
Epoch: 1 / 10
MSE: 3.52e-02; a: 1.3904; b: 1.3368
Epoch: 2 / 10
MSE: 9.21e-03; a: 1.6882; b: 1.1722
Epoch: 3 / 10
MSE: 2.41e-03; a: 1.8406; b: 1.0881
Epoch: 4 / 10
MSE: 6.30e-04; a: 1.9185; b: 1.0450
Epoch: 5 / 10
MSE: 1.65e-04; a: 1.9583; b: 1.0230
Epoch: 6 / 10
MSE: 4.31e-05; a: 1.9787; b: 1.0118
Epoch: 7 / 10
MSE: 1.13e-05; a: 1.9891; b: 1.0060
Epoch: 8 / 10
MSE: 2.95e-06; a: 1.9944; b: 1.0031
Epoch: 9 / 10
MSE: 7.71e-07; a: 1.9971; b: 1.0016
Epoch: 10 / 10
MSE: 2.02e-07; a: 1.9985; b: 1.0008
```



# Linear regression with SGD

## – Comparison of two ways

### The PyTorch way

```
▶ epochs = 10
for epoch in range(epochs):
    for x_batch, y_batch in mydataloader:
        # build computation graph
        yhat = mymodel(x_batch)
        myloss = criterion(y_batch, yhat)

        # optimization
        optimizer.zero_grad() # zero out gradient
        myloss.backward() # back propagation
        optimizer.step() # update parameters
```



Agnostic to datasets, models,  
losses and optimization methods!

### The naive way

```
▶ for epoch in range(nepochs):
    for batch in range(round(n / batch_size)):
        start = batch * batch_size
        end = start + batch_size
        # perform update on a batch
        x_batch = x[start:end]
        y_batch = y[start:end]

        # build computation graph
        y_hat = a * x_batch + b
        myloss = torch.mean((y_batch - y_hat)**2)

        # gradient calculation
        myloss.backward()

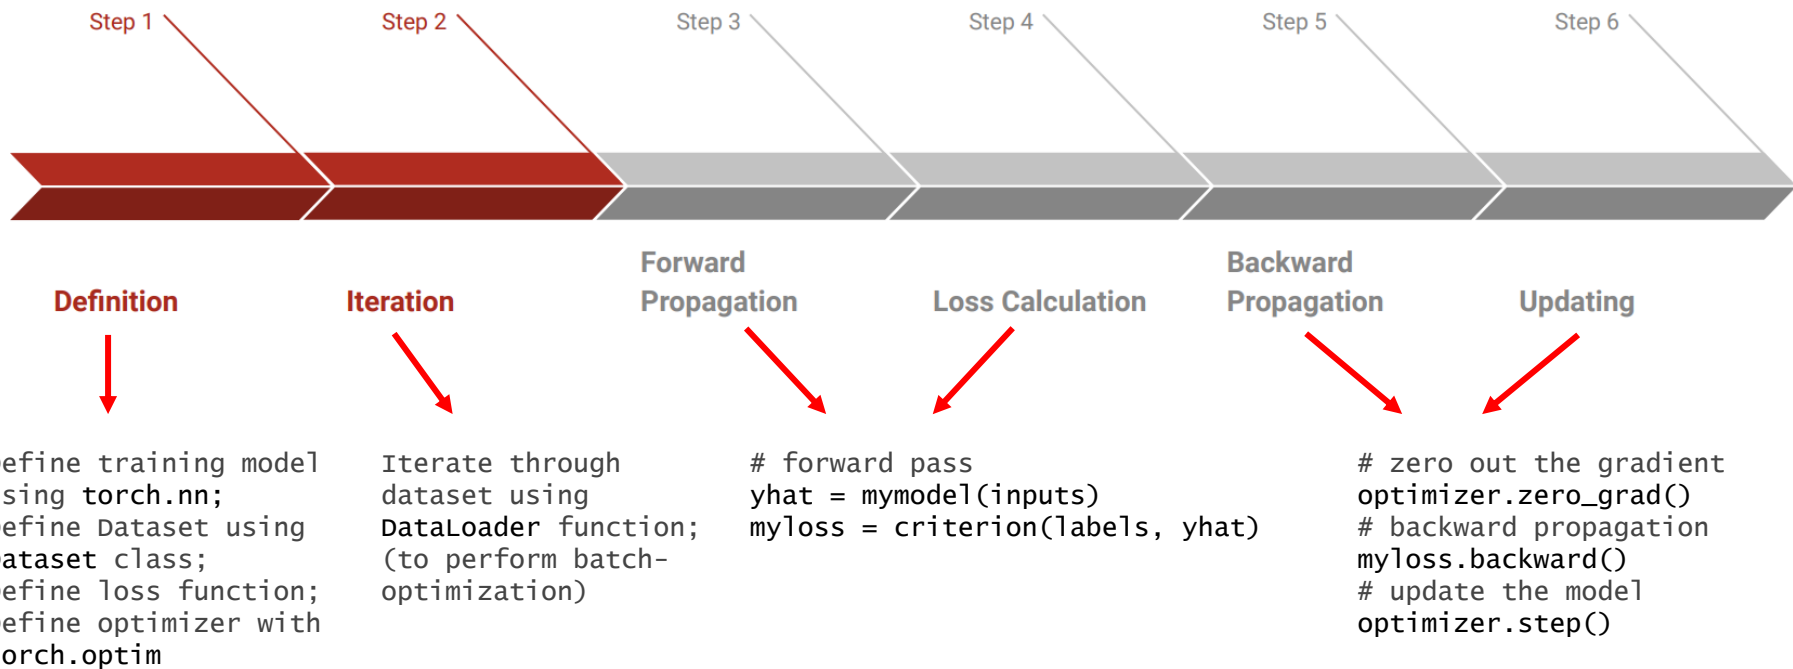
        # SGD update
        with torch.no_grad():
            a -= lr * a.grad
            b -= lr * b.grad

        # avoid gradient accumulation
        a.grad.zero_()
        b.grad.zero_()
```



# ML workflow in PyTorch

*Credit: HPRC Short Course by Jian Tao, TAMU*



# Outline

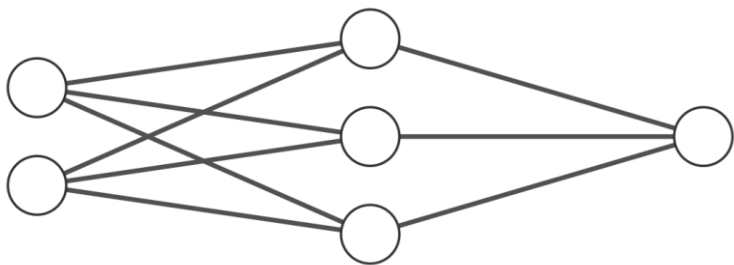
- Set up working environment
- Introduction to PyTorch framework
- Deep learning with PyTorch
  - Regression using ReLU neural net
  - Image classification with CNN
- Additional resources



# From linear model to shallow ReLU network

Usage: `nn.Linear(in_features, n_out_features)`

## Shallow ReLU neural net



Input Layer  $\in \mathbb{R}^2$

Hidden Layer  $\in \mathbb{R}^3$

Output Layer  $\in \mathbb{R}^1$

ReLUUnn2:  $\mathbf{x} \in \mathbb{R}^2 \mapsto \mathbf{w}_2^\top (\mathbf{W}_1 \mathbf{x} + \mathbf{b}_1)_+ + b_2 \in \mathbb{R}$

$\mathbf{W}_1 \in \mathbb{R}^{3 \times 2}, \mathbf{b}_1 \in \mathbb{R}^3, \mathbf{w}_2 \in \mathbb{R}^3, b_2 \in \mathbb{R}$



```
# shallow relu net
```

```
ReLUUnn2 = nn.Sequential(  
    nn.Linear(2, 3),  
    nn.ReLU(),  
    nn.Linear(3, 1)  
)  
for param in ReLUUnn2.parameters():  
    print(type(param.data), param.shape)
```

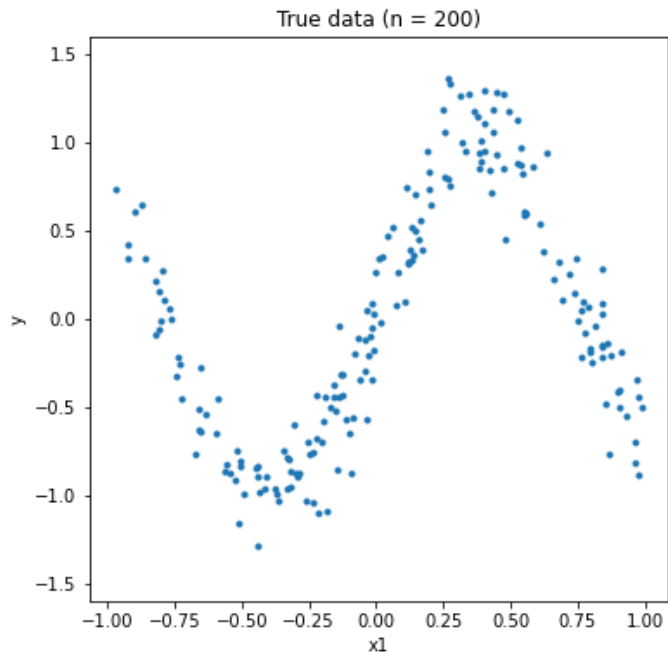


```
<class 'torch.Tensor'> torch.Size([3, 2])  
<class 'torch.Tensor'> torch.Size([3])  
<class 'torch.Tensor'> torch.Size([1, 3])  
<class 'torch.Tensor'> torch.Size([1])
```



# Fit a ReLU neural network to Sine function

What dose the data look like?



$y = 4 \sin(x_1) + \mathcal{N}(0, 1)$  where  $\mathbf{x} \sim \text{Unif}([-1, 1]^5)$

What is our training model?

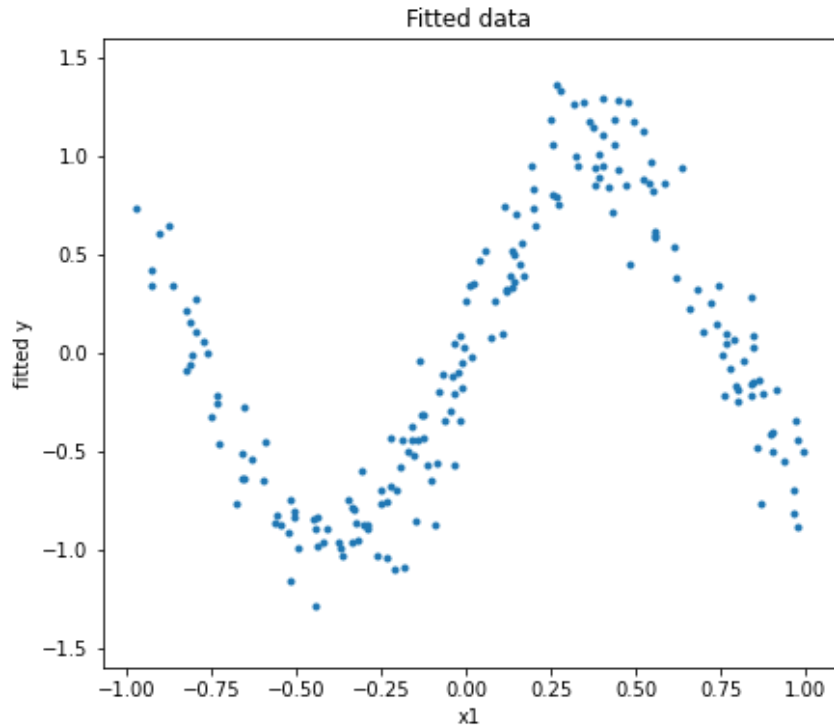
```
# initialize a training model
train_model = nn.Sequential(
    nn.Linear(5, 256),
    nn.ReLU(),
    nn.Linear(256, 128),
    nn.ReLU(),
    nn.Linear(128, 1)
)
train_model
```

```
Sequential(
  (0): Linear(in_features=5, out_features=256, bias=True)
  (1): ReLU()
  (2): Linear(in_features=256, out_features=128, bias=True)
  (3): ReLU()
  (4): Linear(in_features=128, out_features=1, bias=True)
)
```

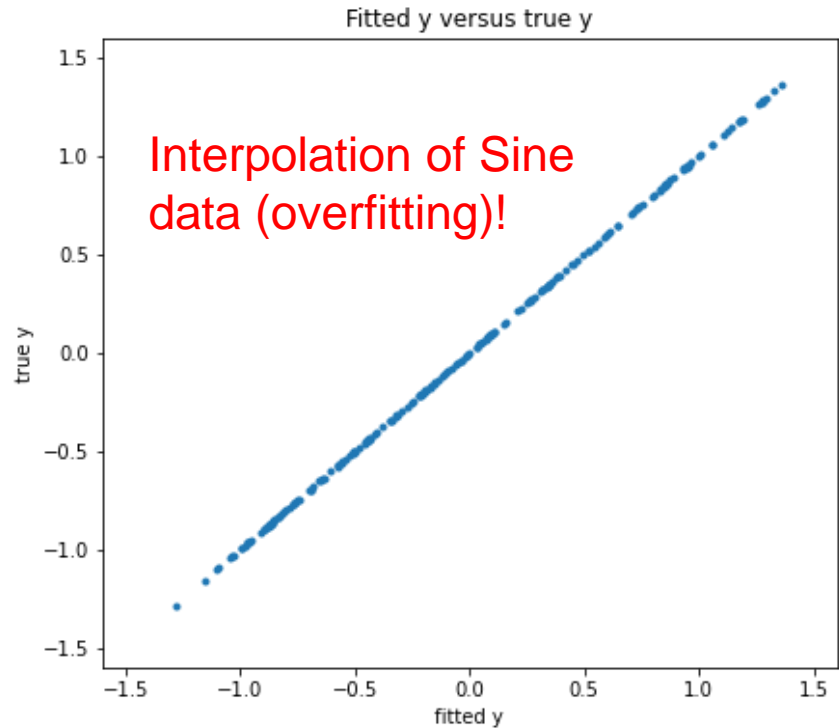
- A two-layer ReLU neural network with 256 hidden units in the first layer, and 128 hidden units in the second layer



# Fit a ReLU neural network to Sine function

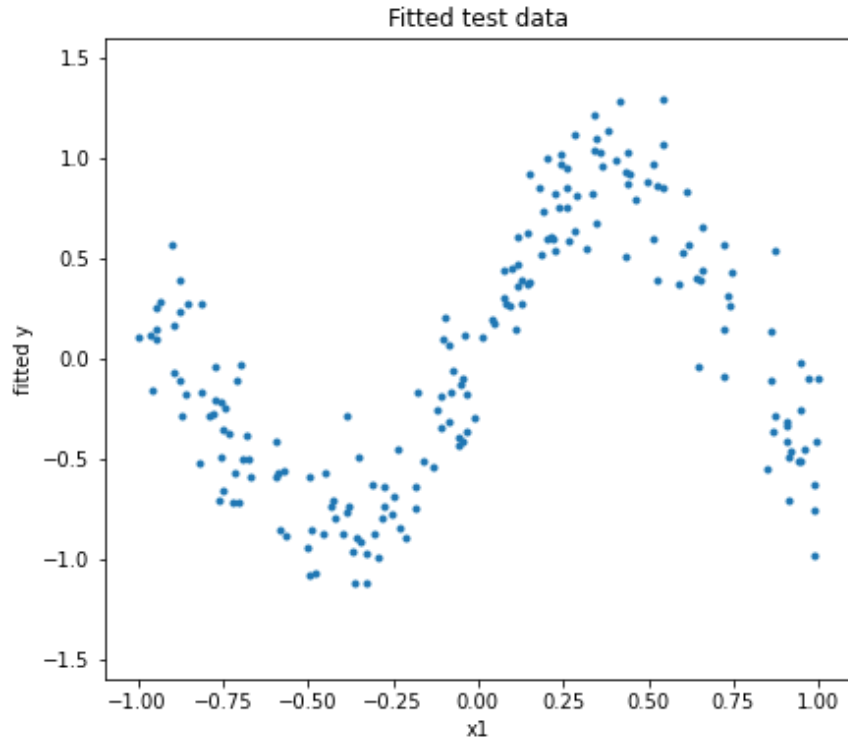


## Trained model on **TRAINING** data

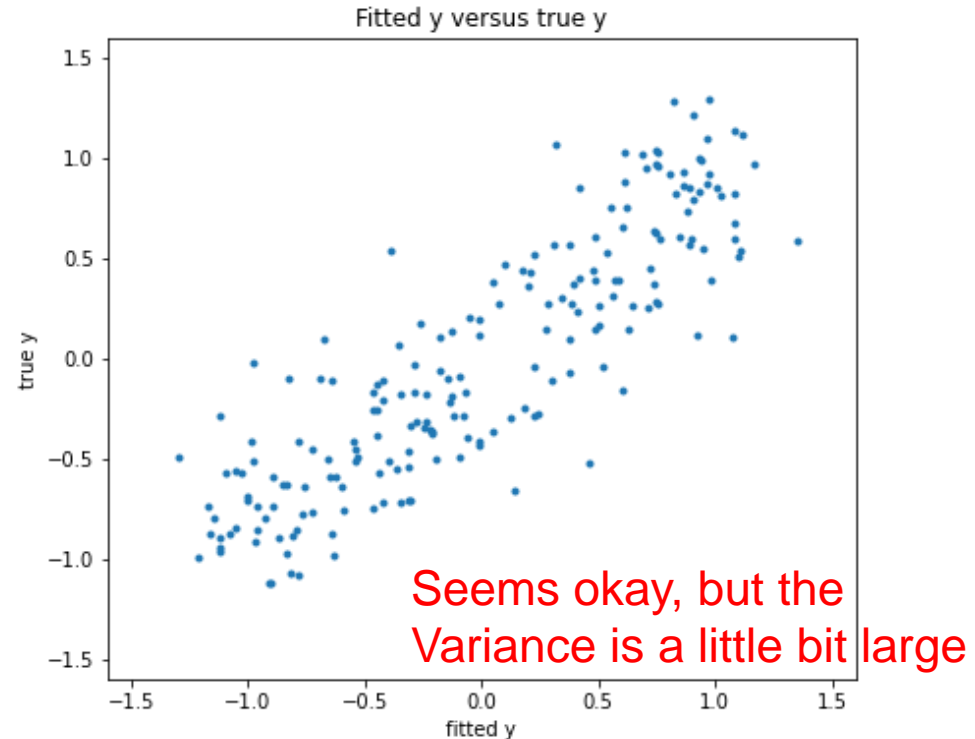




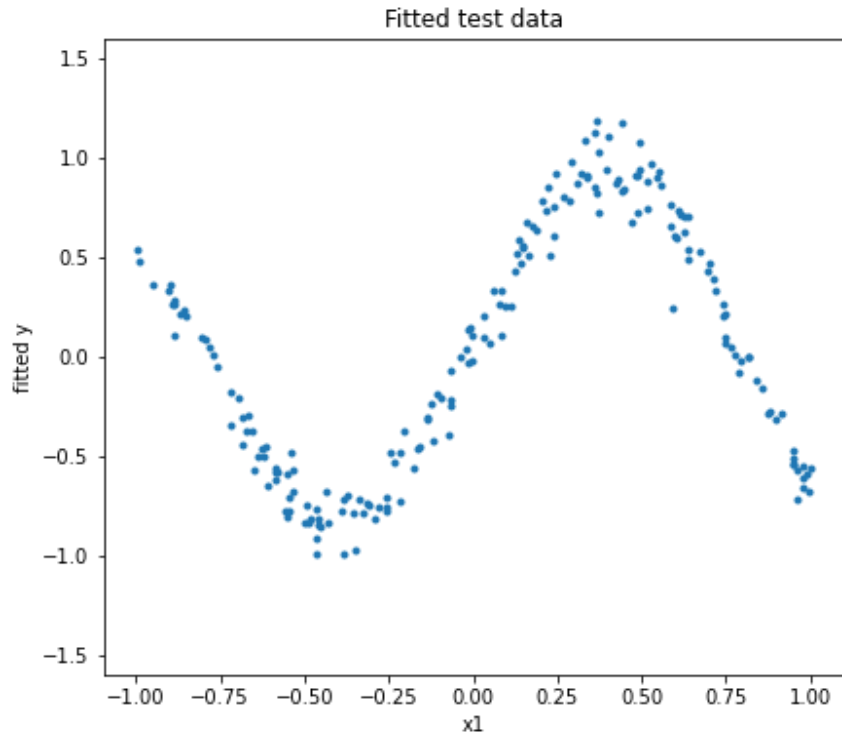
## Fit a ReLU neural network to Sine function



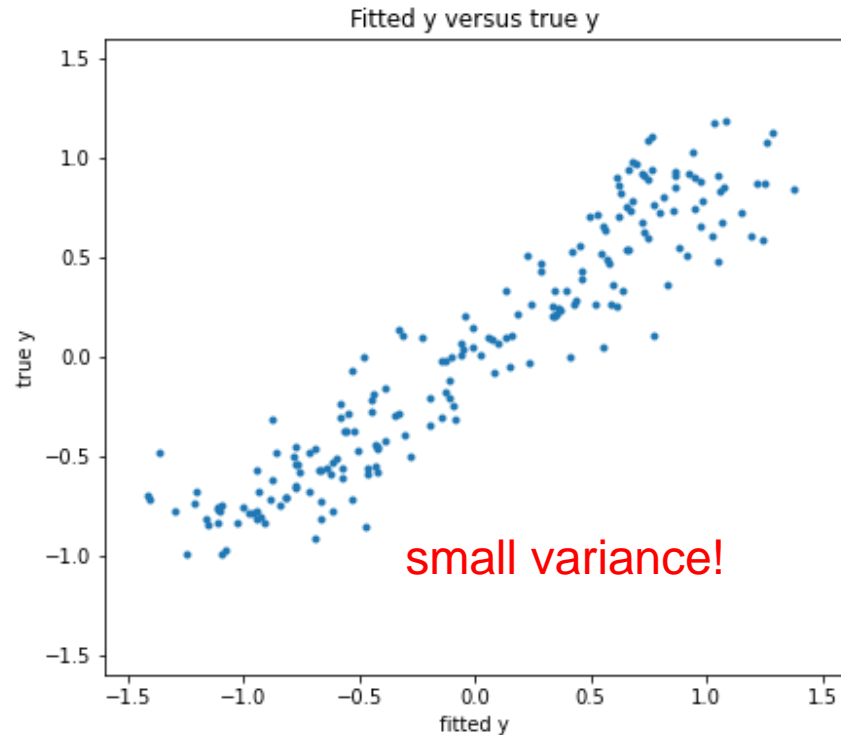
## Trained model on **TESTING** data



## Fit a ReLU neural network to Sine function



## Trained model on TESTING data and with REGULARIZATION



Code: `optimizer = optim.Adam(train_model.parameters(), lr = 0.001, weight_decay = 0.01)`



# Regularization of neural net

## - To achieve better generalization

- Explicit regularization:

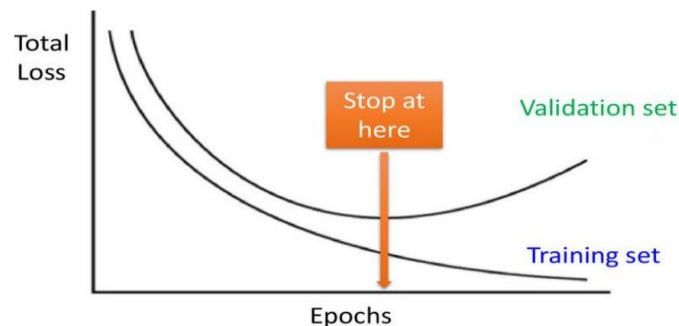
$$\min_{\mathbf{W}} \frac{1}{m} \sum_{i=1}^m \ell(\mathbf{y}_i, \text{DNN}_{\mathbf{W}}(\mathbf{x}_i)) + \lambda \Omega(\mathbf{W})$$

- Norms of weight matrices
- Norms of gradients / Jacobians
- ...

- Miscellaneous: A neural net structure can also induce regularization. E.g. The convolutional neural net (CNN)

- Implicit regularization:

- The regularization that is not built in the objective function (the loss)
- Regularization induced by an **optimization algorithm**: SGD tends to find a solution with small norm (regularized solution)
- **Early stopping**; batch normalization; dropout



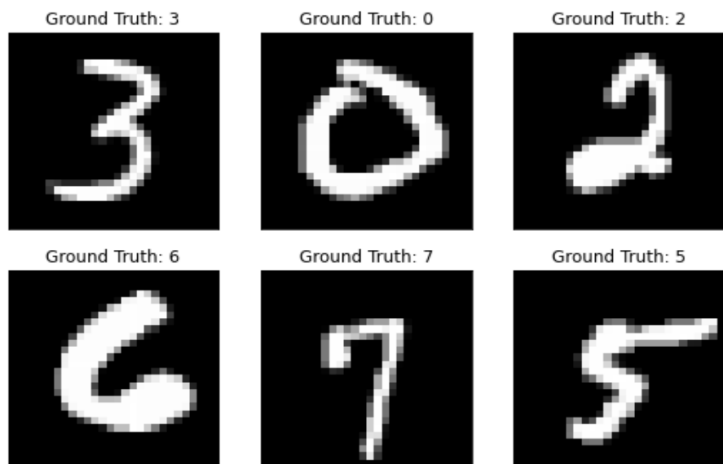
Credit: CSCI 5980/8980: Think Deep Learning offered by Ju Sun, UMN



# Image classification using CNN

## Dataset: MNIST

- One of the benchmark datasets of deep learning for image classification
- Classifying digits as 0, 1, ..., 9

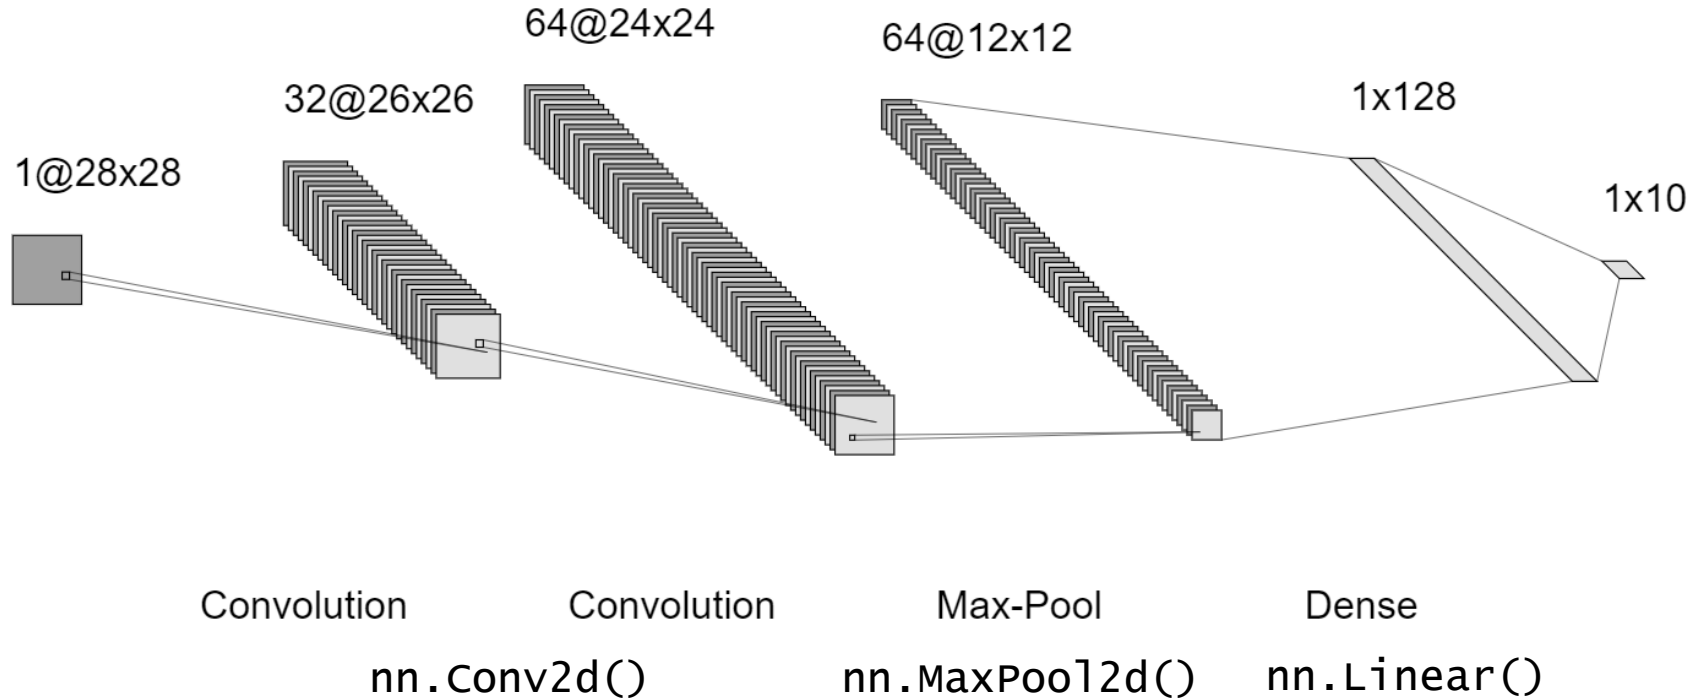


## Our CNN model

```
model = nn.Sequential(  
    nn.Conv2d(1, 32, 3, 1),  
    nn.ReLU(),  
    nn.Conv2d(32, 64, 3, 1),  
    nn.ReLU(),  
    nn.MaxPool2d(2),  
    nn.Dropout(0.25),  
    nn.Flatten(),  
    nn.Linear(9216, 128),  
    nn.ReLU(),  
    nn.Dropout(0.5),  
    nn.Linear(128, 10),  
    nn.LogSoftmax(dim = 1)  
)
```

# Image classification using CNN

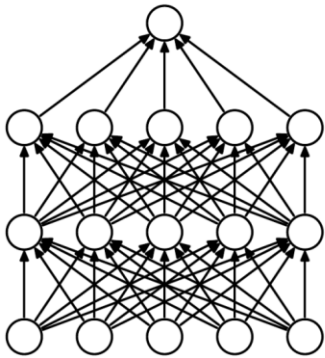
Our CNN structure with activation functions, dropout regularization in between layers.



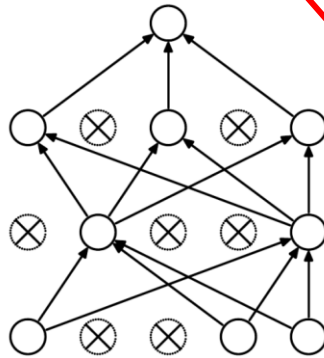
# Image classification using CNN

Dropout for regularization:

- randomly kills inner neurons with some probability  $p$



(a) Standard Neural Net



(b) After applying dropout.

Credit: Srivastava, N., Hinton, G., Krizhevsky, A., Sutskever, I., & Salakhutdinov, R. (2014). Dropout: a simple way to prevent neural networks from overfitting. *The journal of machine learning research*, 15(1), 1929-1958.

## Our CNN model

```
model = nn.Sequential(  
    nn.Conv2d(1, 32, 3, 1),  
    nn.ReLU(),  
    nn.Conv2d(32, 64, 3, 1),  
    nn.ReLU(),  
    nn.MaxPool2d(2),  
    nn.Dropout(0.25),  
    nn.Flatten(),  
    nn.Linear(9216, 128),  
    nn.ReLU(),  
    nn.Dropout(0.5),  
    nn.Linear(128, 10),  
    nn.LogSoftmax(dim = 1)  
)
```



# Image classification using CNN

The total number of trainable parameters is 1,199,882!

Layer (type)	Input Shape	Param #	Tr. Param #
Conv2d-1	[1, 1, 28, 28]	320	320
ReLU-2	[1, 32, 26, 26]	0	0
Conv2d-3	[1, 32, 26, 26]	18,496	18,496
ReLU-4	[1, 64, 24, 24]	0	0
MaxPool2d-5	[1, 64, 24, 24]	0	0
Dropout-6	[1, 64, 12, 12]	0	0
Flatten-7	[1, 64, 12, 12]	0	0
Linear-8	[1, 9216]	1,179,776	1,179,776
ReLU-9	[1, 128]	0	0
Dropout-10	[1, 128]	0	0
Linear-11	[1, 128]	1,290	1,290
LogSoftmax-12	[1, 10]	0	0
Total params: 1,199,882			
Trainable params: 1,199,882			
Non-trainable params: 0			



## Check GPU availability

CPU

```
# check device availability
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
print(device)

cpu
```

GPU

```
# check device availability
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
print(device)

cuda
```





# Download and preprocess the MNIST training / testing dataset

Other transforms:  
cropping, translation,  
rotation, padding...

```
# define image transformation
transform=transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize((0.1307,), (0.3081,))
])

# import MNIST dataset
dataset1 = datasets.MNIST('../data', train = True, download = True, transform = transform)
dataset2 = datasets.MNIST('../data', train = False, transform = transform)

# only use 2000 training images and 20 testing images
train_index = np.random.choice(range(60000), 2000, replace = False)
test_index = np.random.choice(range(10000), 20, replace = False)
train_data = Subset(dataset1, train_index)
test_data = Subset(dataset2, test_index)
```



## Download and preprocess the MNIST training / testing dataset

Most popular benchmark datasets can be loaded via [torchvision](#) library

```
# define image transformation
transform=transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize((0.1307,), (0.3081,))
])

# import MNIST dataset
dataset1 = datasets.MNIST('../data', train = True, download = True, transform = transform)
dataset2 = datasets.MNIST('../data', train = False, transform = transform)

# only use 2000 training images and 20 testing images
train_index = np.random.choice(range(60000), 2000, replace = False)
test_index = np.random.choice(range(10000), 20, replace = False)
train_data = Subset(dataset1, train_index)
test_data = Subset(dataset2, test_index)
```

and don't forget to prepare DataLoader for training



## Define a training utils function

```
# a utils function for training
def train(model, device, train_loader, optimizer, epoch):
    model.train() # enable dropout
    correct = 0
    for batch_idx, (data, target) in enumerate(train_loader):
        # transfer batches of data to specified device
        data, target = data.to(device), target.to(device)

        # build computation graph
        output = model(data)
        loss = F.nll_loss(output, target)

        # the optimization part
        optimizer.zero_grad()
        loss.backward()
        optimizer.step()

    # count the corrected classified cases
    pred = output.argmax(dim = 1, keepdim = True) # get the index of the max log-probability
    correct += pred.eq(target.view_as(pred)).sum().item()

    print(f"Epoch {epoch + 1}:")
    print(f"Training accuracy: {100. * correct / len(train_loader.dataset): .2f} %")
```

Similar as before,  
nothing new



# Train our CNN



```
# the training
model = model.to(device)
optimizer = optim.Adam(model.parameters(), lr = 0.005)

nepochs = 10
for epoch in range(nepochs):
    train(model, device, train_loader, optimizer, epoch)
    test(model, device, test_loader)
```

```
↳ Epoch 1: Training accuracy: 41.10 %; Testing accuracy: 154/200 (77%)
Epoch 2: Training accuracy: 78.35 %; Testing accuracy: 179/200 (90%)
Epoch 3: Training accuracy: 88.25 %; Testing accuracy: 186/200 (93%)
Epoch 4: Training accuracy: 92.25 %; Testing accuracy: 185/200 (92%)
Epoch 5: Training accuracy: 94.15 %; Testing accuracy: 195/200 (98%)
Epoch 6: Training accuracy: 96.05 %; Testing accuracy: 191/200 (96%)
Epoch 7: Training accuracy: 97.00 %; Testing accuracy: 192/200 (96%)
Epoch 8: Training accuracy: 96.55 %; Testing accuracy: 193/200 (96%)
Epoch 9: Training accuracy: 96.85 %; Testing accuracy: 194/200 (97%)
Epoch 10: Training accuracy: 97.60 %; Testing accuracy: 195/200 (98%)
```



## Save / load the trained model

- State dictionary method: only save parameter values:



```
# save parameter values
torch.save(model.state_dict(), "model_dict.pt")

# When one wants to use saved model next time:
# 1. initialize the same CNN model
model_new = myCNN()
# 2. load parameter values
model_new.load_state_dict(torch.load("model_dict.pt"))
model_new.eval()
```



# Save / load the trained model

- Checkpoint method: save for resuming training later

```
# save checkpoint for resuming training
torch.save({
    'epoch': epoch,
    'model_state_dict': model.state_dict(),
    'optimizer_state_dict': optimizer.state_dict()
}, "model_checkpoint.pt")

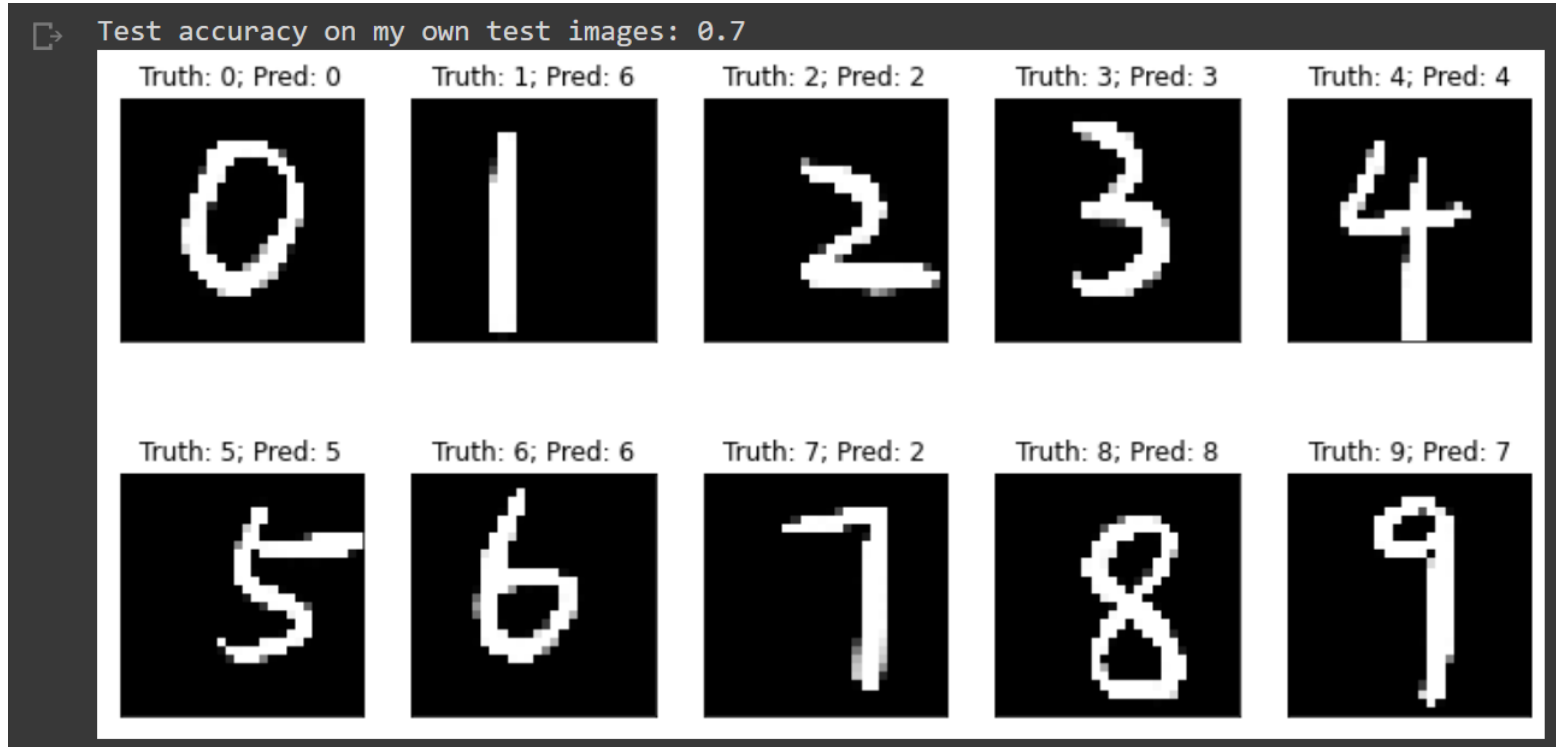
# When one wants to use saved model next time:
# 1. initialize class and optimizer
model_new = myCNN()
optimizer_new = optim.Adam(model_new.parameters(), lr = 0.005)
# 2. load checkpoints for both model and optimizer
checkpoint = torch.load("model_checkpoint.pt")
model_new.load_state_dict(checkpoint['model_state_dict'])
optimizer_new.load_state_dict(checkpoint['optimizer_state_dict'])
epoch = checkpoint['epoch']
# 3. do whatever you like :)
# - either - evaluate the mode and do some inference
model_new.eval()
# - or - resume training from your last checkpoint
model_new.train()
```

Check [saving and loading models](#) for more details



# Robustness issue of CNN

Test accuracy only 0.7.  
Solution: data augmentation!



# Robustness issue of CNN

Solution: adversarial training!

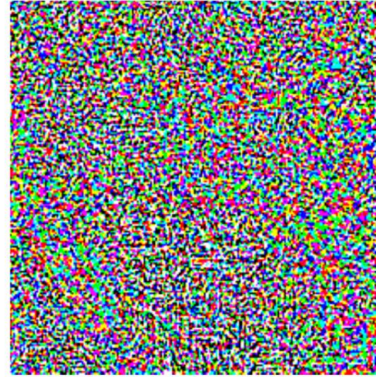


$x$

“panda”

57.7% confidence

$+ .007 \times$



$\text{sign}(\nabla_x J(\theta, x, y))$

“nematode”

8.2% confidence

$=$



$x + \epsilon \text{sign}(\nabla_x J(\theta, x, y))$

“gibbon”

99.3 % confidence

*Credit: Goodfellow, I. J., Shlens, J., & Szegedy, C. (2014). Explaining and harnessing adversarial examples. arXiv preprint arXiv:1412.6572.*





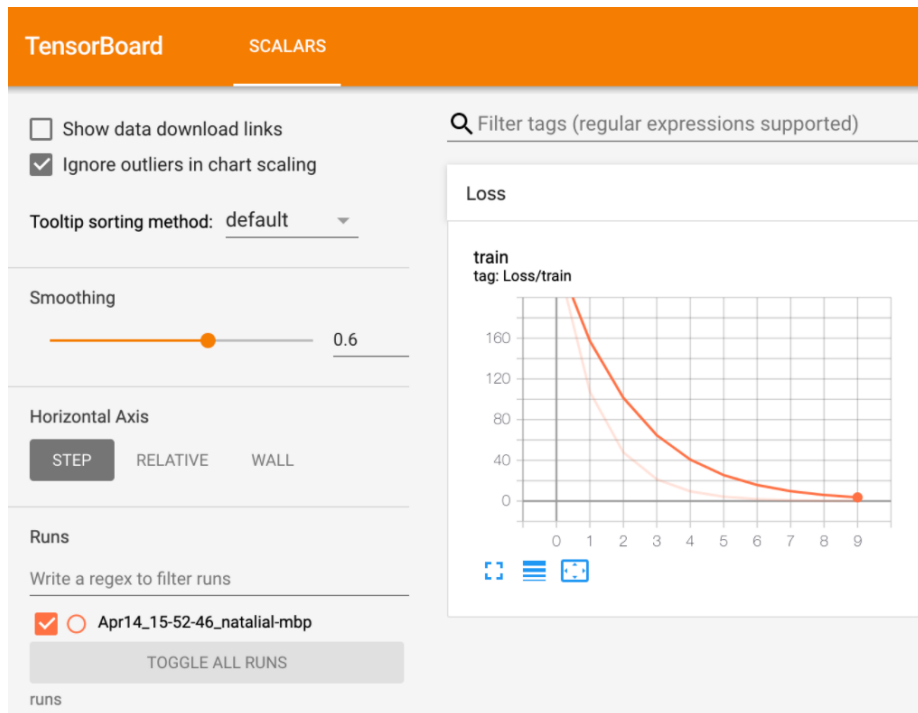
# Outline

- Set up working environment
- Introduction to PyTorch framework
- Deep learning with PyTorch
- Other resources



# Visualization with TensorBoard

- Track and visualize:
  - Metrics such as loss and accuracy
  - The model graph
  - Histograms, images and much more
- Some resources:
  - [How to use TensorBoard with PyTorch](#)
  - [Visualizing models, data, and training with TensorBoard](#)
  - [PyTorch TensorBoard support](#)



# Tutorials and courses:

## Books and tutorials:

- [Dive into Deep Learning](#) (livebook)
- [Deep Learning](#) by Ian Goodfellow
- [Official PyTorch tutorial](#)
- [Deep learning with Python](#) (livebook)
- [UvA DL Notebooks](#)

## Courses:

- [DL/ML tutorial](#) by Hung-Yi Lee
- [Deep learning](#) course by Yann LeCun
- [Deep learning with Pytorch](#)
- [Stanford STAT385 series](#)
- [Think Deep Learning](#) by Ju Sun





UNIVERSITY OF MINNESOTA

**Driven to Discover®**

Crookston Duluth Morris Rochester Twin Cities

The University of Minnesota is an equal opportunity educator and employer.