

Introduction to PyTorch

Yifei Liu

March, 2022



UNIVERSITY OF MINNESOTA

Driven to Discover®

Outline

- Set up working environment
 - Coding platform
 - Create conda environment
- Introduction to PyTorch framework
- Deep learning with PyTorch
- Additional resources



Coding platform

Jupyter Notebook / Lab

On your own computer:

- Open [Anaconda](#) Prompt
- Run `jupyter notebook`

On HPC resources from MSI:

- Need to be in a research group
- [UMN VPN](#) (off-campus only)
- Visit notebooks.msi.umn.edu
- Choose and start a server

Alternative choice: [Colab Notebook](#)

VS Code

- Install the [Python extension](#)
- Check [here](#) for:
 - Path configuration
 - Autocompletion and hints
 - Debugging ...

To use VS Code from a remote host:

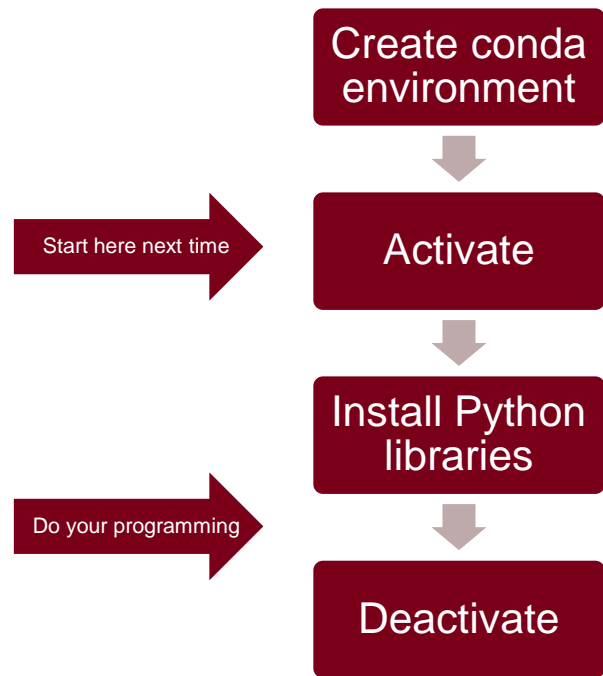
- [Remote development extension](#)

Other choices:

- PyCharm, Spyder, Sublime ...



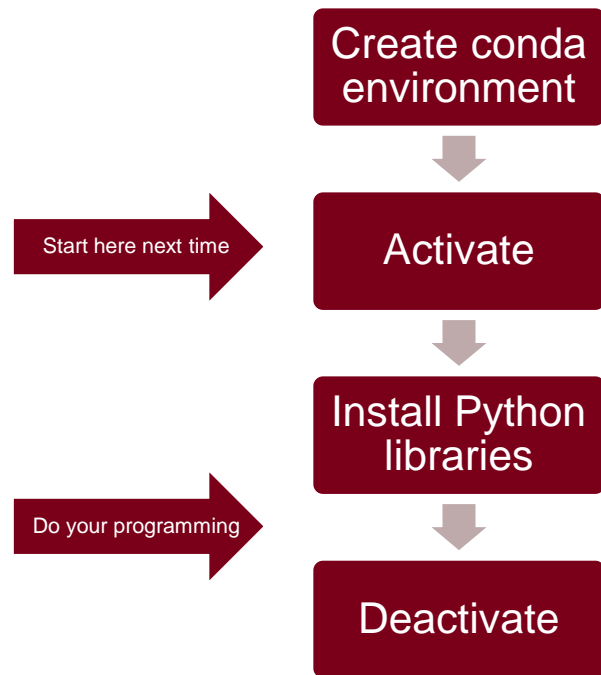
Create a conda environment for your Python project



- An alternative way to Virtualenv
- Manage dependencies (**different versions** of Python and some libraries) **within an isolated project**
 - An example: Python library [foolbox](#)
- Key commands:
 - `conda create -n <envname> python=3.8.3`
 - `conda activate <envname>`
 - `conda install <package>=<version>`
 - `conda deactivate`
- Check [Conda Cheat Sheet](#)



A hands-on exercise



- Log in to notebooks.msi.umn.edu
- We want to create a conda environment called “8056ex”:
 - Python (version 3.8.3)
 - NumPy (version 1.18.1)
- Some commands
 - `conda create -n 8056ex python=3.8.3`
 - `conda activate 8056ex`
 - `conda install numpy=1.18.1`
- Specific to Jupyter Notebook / Lab:
 - `conda install ipykernel`
 - `python -m ipykernel install --user --name 8056ex --display-name “8056ex”`



Installation of PyTorch

- Install PyTorch from <http://pytorch.org/>

PyTorch Build	Stable (1.11.0)	Preview (Nightly)	LTS (1.8.2)	
Your OS	Linux	Mac	Windows	
Package	Conda	Pip	LibTorch	Source
Language	Python	C++ / Java		
Compute Platform	CUDA 10.2	CUDA 11.3	ROCm 4.2 (beta)	CPU
Run this Command:	<code>conda install pytorch torchvision torchaudio cudatoolkit=11.3 -c pytorch</code>			
Previous versions of PyTorch >				

Run this within your conda environment



UMN computing resources

MSI: <https://www.msi.umn.edu/>

Introduction to Minnesota Supercomputing Institute (MSI)	This tutorial is geared to new MSI users and will provide a highlevel introduction to the facilities and computational resources at MSI.	02/01/2022
Introduction to Linux	This tutorial will provide an introduction to the Linux operating system, with particular attention paid to working from the command line	02/03/2022
Programming with Python	Introduction to fundamentals of programming using the python language.	02/08/2022
Job Submission and Scheduling at MSI	This tutorial will introduce users to MSI supercomputers, and provide an overview of how to submit calculations to the job schedulers	02/10/2022
Interactive Computing at MSI	This two part tutorial will introduce you to the concept of interactive high performance computing, and provide attendees hands-on experience running interactive parallel jobs on the Mesabi HPC	02/15/2022
Data Storage Systems and Data Analysis Workflows for Research	In this tutorial you will learn about the data storage systems available for academic research at the University of Minnesota	02/17/2022
Compiling and Debugging at MSI	This tutorial will help users learn the basics of compiling and debugging their code on MSI systems	02/24/2022
Python for Scientific Computing	This session includes efficient data processing with NumPy and Scipy, data visualization, and techniques for using python to drive parallel supercomputing tasks.	03/01/2022
RNA-Seq Analysis	This tutorial covers the basics of differential expression analysis and touches on other RNA-seq topics such as transcriptome assembly.	03/03/2022
Parallel Computing On Agate	This tutorial will help users learn the basics of parallel computation methods, including strategies for collecting calculations together for parallel execution.	03/31/2022

LATIS: <https://latisresearch.umn.edu/>

Feb. 18th 10:00am-noon	Creating Publication Worthy Visualizations without Code	Registration
Feb. 25th 10:00am-noon	Introduction to Computational Text Analysis	Registration
Mar. 4th 10:00am-noon	Reproducible research practices in Excel (yes, Excel)	Registration
Mar. 9th 10:00am-noon	Data Management in transition: Strategies for when you graduate	Registration
Mar. 18th 10:00am-noon [RESCHEDULED]	Advanced Nvivo	Registration
Mar. 25th 10:00am-noon	Introduction to parallel computing	Registration
April 1st 10:00am-noon	Introduction to SQL and Research Databases	Registration



Outline

- Set up working environment
- Introduction to PyTorch framework
 - Tensor, gradient and computation graph
 - Tensor manipulation on GPU
 - Use PyTorch as a general ML framework
- Deep learning with PyTorch
- Additional resources



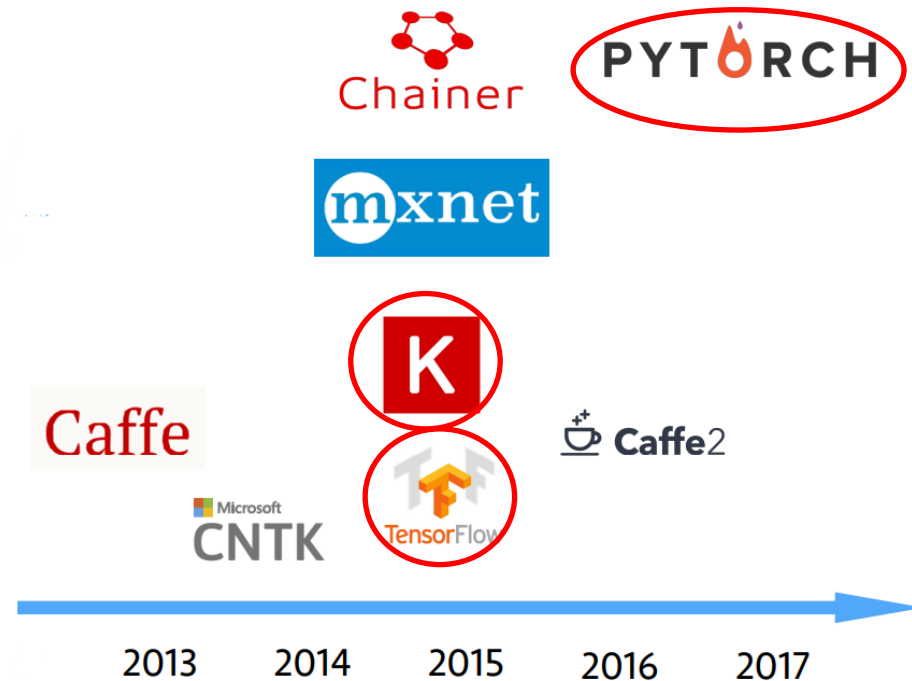
Overview of PyTorch

A fancy version of NumPy but can:

- Trace gradient via autograd
- Accelerate via GPU
- Accommodate a bunch of machine learning functionalities including but not limited to neural network

More importantly,

- Easy interface compared to others
- Easy to debug and understand codes
- Check this [Google Trends](#)



Credit: "Gluon: new MXNet interface to accelerate research"



Tensor

What is a tensor?

- A multidimensional array
- Some examples:
 - 0D-tensor: scalar
 - 1D-tensor: vector
 - 2D-tensor: matrix; non-RGB image
 - 3D-tensor: RGB image (3 channels)
 - 4D-tensor: **one** RGB video clip
 - 5D-tensor: **a collection of** RGB video clips
 - 6D-tensor: ???
- Tensor in PyTorch:
 - Just like arrays in Numpy
 - `t.numpy()` and `torch.from_numpy(a)`

np.array versus torch.tensor

- Similar object creations:
 - All-ones, all-zeros, identity matrix, random...
 - Check [PyTorch tensor creations](#)
- Similar math operations:
 - Indexing, slicing, reshape, transpose, tensor product, element-wise operation...
 - Check [PyTorch tensor math operations](#)

```
import numpy as np
myarray = np.ones(3)
print(myarray)
print(myarray + 1.5)
```

```
[1.  1.  1.]
[2.5  2.5  2.5]
```

```
import torch
mytensor = torch.ones(3)
print(mytensor)
print(mytensor + 1.5)
```

```
tensor([1., 1., 1.])
tensor([2.5000, 2.5000, 2.5000])
```



Tensor (cont.)

- Enable gradient tracing:
 - Upon any tensor creation, set `requires_grad = True`
 - For any existent tensor `t`, call `t.requires_grad_(True)`
- Calculate gradient:
 - Do some calculation from `t` and get `f`
 - Call `f.backward()` for calculation
 - Call `t.grad` to access the gradient
- Disable gradient tracing:
 - (Permanently) `t = t.detach()`
 - (Temporarily) with `torch.no_grad()`:



```
# create tensors
c = torch.tensor(1.)
t = torch.tensor(3., requires_grad = True)
print(c)
print(t)

# do your calculation based on t
s = t**2
f = 3 * s + c

# calculate the gradient w.r.t. t
f.backward()

# access the gradient
t.grad
```

```
tensor(1.)
tensor(3., requires_grad=True)
tensor(18.)
```



Dynamic computation graph

t

c

enable
gradient

do some
calculation

call
.backward()

access the
gradient

create tensors

```
c = torch.tensor(1.)
```

```
t = torch.tensor(3., requires_grad = True)
```

```
print(c)
```

```
print(t)
```

do your calculation based on t

```
s = t**2
```

```
f = 3 * s + c
```

calculate the gradient w.r.t. t

```
f.backward()
```

access the gradient

```
t.grad
```





```
tensor(1.)
```

```
tensor(3., requires_grad=True)
```

```
tensor(18.)
```



Dynamic computation graph



```
# create tensors
c = torch.tensor(1.)
t = torch.tensor(3., requires_grad = True)
print(c)
print(t)

# do your calculation based on t
s = t**2
f = 3 * s + c

# calculate the gradient w.r.t. t
f.backward()

# access the gradient
t.grad
```

tensor(1.)
tensor(3., requires_grad=True)
tensor(18.)




Dynamic computation graph

enable
gradient

do some
calculation

call
.backward()

access the
gradient



```
# create tensors
c = torch.tensor(1.)
t = torch.tensor(3., requires_grad = True)
print(c)
print(t)
```

```
# do your calculation based on t
```

```
s = t**2
```

```
f = 3 * s + c
```

```
# calculate the gradient w.r.t. t
```

```
f.backward()
```

```
# access the gradient
```

```
t.grad
```



```
tensor(1.)
tensor(3., requires_grad=True)
tensor(18.)
```



Dynamic computation graph

enable
gradient

do some
calculation

call
.backward()

access the
gradient

The graph structure is destroyed
Once you call `.backward()`

t

s

c

f



```
# create tensors
c = torch.tensor(1.)
t = torch.tensor(3., requires_grad = True)
print(c)
print(t)

# do your calculation based on t
s = t**2
f = 3 * s + c

# calculate the gradient w.r.t. t
f.backward()


# access the gradient
t.grad
```



```
tensor(1.)
tensor(3., requires_grad=True)
tensor(18.)
```



Dynamic computation graph



```
# create tensors
c = torch.tensor(1.)
t = torch.tensor(3., requires_grad = True)
print(c)
print(t)

# do your calculation based on t
s = t**2
f = 3 * s + c

# calculate the gradient w.r.t. t
f.backward()

# access the gradient
t.grad
```

```
tensor(1.)
tensor(3., requires_grad=True)
tensor(18.)
```



Why should we understand PyTorch computation graph?

- Essentially, PyTorch is using **chain rule** to calculate the gradient
- A computation graph defines how the chain rule applies to your calculation
- **Some potential MISTAKES:**
 - Retrieve gradient of non-leaf nodes
 - Call `.backward()` when there is no graph
 - Gradient accumulation. To solve this issue, call `t.grad.zero_()` before building graph

```
# first back propagation
t = torch.tensor(3., requires_grad = True)
s = t**2
s.backward()
print(t.grad)

# second propagation (WITHOUT RE-CREATING t)
f = 5 * t + 1
f.backward()
print(t.grad)
```

```
tensor(6.)
tensor(11.)
```

should be 5


```
t.grad.zero_() # zero out the gradient
f = 5 * t + 1
f.backward()
print(t.grad)

tensor(5.)
```



Calculation on GPU

- Neural network dated back to 1950s, but became popular **only recently due to huge improvements in computation**



Credit: CSCI 5980/8980: Think Deep Learning offered by Ju Sun, UMN

- GPU allows faster **large-scale matrix multiplication** (or tensor product)
- Use GPU on PyTorch:
 - `torch.cuda.is_available()` # GPU?
 - `t = t.cuda()` # move to GPU
 - `t = t.cpu()` # move to CPU
 - `t.device` # cpu or cuda: index
 - The transmission time can be long
 - Make sure all tensor manipulations are performed on the same device
- Advanced: [train on multiple GPUs](#)



Time comparison: matrix multiplication on CPU and GPU

- We want to perform matrix multiplication $C = AB$ where both A and B are 3000 x 3000 matrices
- Enable GPU on [Google Colab](#) (free!)
 1. Runtime
 2. Change runtime type
 3. Hardware Accelerator: GPU
 4. Runtime -> Restart Runtime

```
d = 3000

# C = A B with numpy on CPU
A = np.random.rand(d, d)
B = np.random.rand(d, d)
begin = time.time()
C = A.dot(B)
print(f"CPU time (numpy): {time.time() - begin} s")

# C = A B with torch on CPU
A_t_cpu = torch.tensor(A)
B_t_cpu = torch.tensor(B)
begin = time.time()
C = torch.mm(A_t_cpu, B_t_cpu)
print(f"CPU time (torch): {time.time() - begin} s")

# C = A B with torch on GPU
# make sure GPU is available
A_t_gpu = torch.tensor(A).cuda()
B_t_gpu = torch.tensor(B).cuda()
begin = time.time()
C = torch.mm(A_t_gpu, B_t_gpu)
print(f"GPU time (torch): {time.time() - begin} s")
```

```
CPU time (numpy): 1.604166030883789 s
CPU time (torch): 1.5383508205413818 s
GPU time (torch): 0.0011920928955078125 s
```



GPU not always better

- Calculation on GPU should be optimized smartly since
 - GPU has **limited memory** compared to CPU
 - **Transmission time** from CPU to GPU can be large

```
▶ d = 3000
begin = time.time()
A = torch.rand((d, d)).cuda()
B = torch.rand((d, d)).cuda()
print(f"Time from CPU to GPU: {time.time() - begin} s")
begin = time.time()
C = torch.mm(A, B)
print(f"GPU time (torch): {time.time() - begin} s")

☞ Time from CPU to GPU: 0.15282464027404785 s
GPU time (torch): 0.0008263587951660156 s
```

NVIDIA-SMI 440.100			Driver Version: 440.100			CUDA Version: 10.2		
GPU	Name	Persistence-M	Bus-Id	Disp.A	Volatile Uncorr. ECC			
Fan	Temp	Perf	Pwr:Usage/Cap	Memory-Usage	GPU-Util	Compute	M.	
0	TITAN RTX	Off	00000000:1A:00.0	Off	N/A			
41%	56C	P2	129W / 280W	13556MiB / 24220MiB	61%	Default		
1	TITAN RTX	Off	00000000:1B:00.0	Off	N/A			
41%	55C	P2	127W / 280W	18798MiB / 24220MiB	60%	Default		
2	TITAN RTX	Off	00000000:3D:00.0	Off	N/A			
41%	48C	P2	121W / 280W	13616MiB / 24220MiB	62%	Default		
3	TITAN RTX	Off	00000000:3E:00.0	Off	N/A			
41%	56C	P2	121W / 280W	20962MiB / 24220MiB	61%	Default		
4	TITAN RTX	Off	00000000:88:00.0	Off	N/A			
41%	53C	P2	162W / 280W	13590MiB / 24220MiB	92%	Default		
5	TITAN RTX	Off	00000000:89:00.0	Off	N/A			
41%	59C	P2	155W / 280W	18832MiB / 24220MiB	61%	Default		
6	TITAN RTX	Off	00000000:B1:00.0	Off	N/A			
41%	55C	P2	169W / 280W	13650MiB / 24220MiB	90%	Default		
7	TITAN RTX	Off	00000000:B2:00.0	Off	N/A			
41%	56C	P2	109W / 280W	20996MiB / 24220MiB	91%	Default		



ML in PyTorch

- PyTorch provides a framework for general ML algorithm **including but not limited to deep learning**
- How does a ML algorithm work?
 - Prepare your dataset
 - Choose a training model
 - Choose a loss function
 - **Optimization** (calculation; gradient)
 - **Evaluation** and inference
- The go-to optimization: SGD
 - Instead of using gradient calculated from all training samples, we **only use the gradient from a randomly chosen sample**
- Practical one: (Mini-) Batch SGD
 - Use gradient from **a batch of samples**
 - **Batch size**: # samples in one batch
 - **Epoch**: **a full pass** of all training samples
 - Special cases:
 - Batch size = 1: vanilla SGD
 - Batch size = n: GD



Linear regression with SGD

– The naive way

- The true model: $y = 2 * x + 1$
- Some other setups:
 - $n = 2000$, $x \sim \text{Uniform}(0, 1)$
 - Number of epochs: 10
 - Batch size: 200 (10 batches per epoch)
 - Learning rate: 0.05
 - Initialize both intercept and slope with $\text{Uniform}(0, 1)$



```
n = 2000
x = torch.rand(n)

# create dataset with true model
a0 = 2
b0 = 1
y = a0 * x + b0

# set up optimization parameter of SGD
a = torch.rand(1, requires_grad = True)
b = torch.rand(1, requires_grad = True)
nepochs = 10
batch_size = 200
lr = 0.5
```



```
▶ for epoch in range(nepochs):  
    for batch in range(round(n / batch_size)):  
        start = batch * batch_size  
        end = start + batch_size  
        # perform update on a batch  
        x_batch = x[start:end]  
        y_batch = y[start:end]  
  
        # build computation graph  
        y_hat = a * x_batch + b  
        myloss = torch.mean((y_batch - y_hat)**2)  
  
        # gradient calculation  
        myloss.backward()  
  
        # SGD update  
        with torch.no_grad():  
            a -= lr * a.grad  
            b -= lr * b.grad  
  
        # avoid gradient accumulation  
        a.grad.zero_()  
        b.grad.zero_()  
  
    print(f"Epoch: {epoch + 1} / {nepochs}")  
    print(f"MSE: {myloss: .2e}; a: {a.item():.4f}; b: {b.item():.4f}")
```

Outer iteration of SGD (epoch)

Linear regression with SGD – The naive way

Output:

```
Epoch: 1 / 10  
MSE: 3.69e-02; a: 1.4460; b: 1.2822  
Epoch: 2 / 10  
MSE: 9.24e-03; a: 1.7226; b: 1.1413  
Epoch: 3 / 10  
MSE: 2.32e-03; a: 1.8611; b: 1.0707  
Epoch: 4 / 10  
MSE: 5.81e-04; a: 1.9305; b: 1.0354  
Epoch: 5 / 10  
MSE: 1.46e-04; a: 1.9652; b: 1.0177  
Epoch: 6 / 10  
MSE: 3.65e-05; a: 1.9826; b: 1.0089  
Epoch: 7 / 10  
MSE: 9.15e-06; a: 1.9913; b: 1.0044  
Epoch: 8 / 10  
MSE: 2.29e-06; a: 1.9956; b: 1.0022  
Epoch: 9 / 10  
MSE: 5.75e-07; a: 1.9978; b: 1.0011  
Epoch: 10 / 10  
MSE: 1.44e-07; a: 1.9989; b: 1.0006
```





```
for epoch in range(nepochs):  
    for batch in range(round(n / batch_size)):  
        start = batch * batch_size  
        end = start + batch_size  
        # perform update on a batch  
        x_batch = x[start:end]  
        y_batch = y[start:end]  
  
        # build computation graph  
        y_hat = a * x_batch + b  
        myloss = torch.mean((y_batch - y_hat)**2)  
  
        # gradient calculation  
        myloss.backward()  
  
        # SGD update  
        with torch.no_grad():  
            a -= lr * a.grad  
            b -= lr * b.grad  
  
        # avoid gradient accumulation  
        a.grad.zero_()  
        b.grad.zero_()  
  
print(f"Epoch: {epoch + 1} / {nepochs}")  
print(f"MSE: {myloss: .2e}; a: {a.item():.4f}; b: {b.item():.4f}")
```

Inner iteration
of SGD (batch)

Linear regression with SGD – The naive way

Output:

```
Epoch: 1 / 10  
MSE: 3.69e-02; a: 1.4460; b: 1.2822  
Epoch: 2 / 10  
MSE: 9.24e-03; a: 1.7226; b: 1.1413  
Epoch: 3 / 10  
MSE: 2.32e-03; a: 1.8611; b: 1.0707  
Epoch: 4 / 10  
MSE: 5.81e-04; a: 1.9305; b: 1.0354  
Epoch: 5 / 10  
MSE: 1.46e-04; a: 1.9652; b: 1.0177  
Epoch: 6 / 10  
MSE: 3.65e-05; a: 1.9826; b: 1.0089  
Epoch: 7 / 10  
MSE: 9.15e-06; a: 1.9913; b: 1.0044  
Epoch: 8 / 10  
MSE: 2.29e-06; a: 1.9956; b: 1.0022  
Epoch: 9 / 10  
MSE: 5.75e-07; a: 1.9978; b: 1.0011  
Epoch: 10 / 10  
MSE: 1.44e-07; a: 1.9989; b: 1.0006
```





```
for epoch in range(nepochs):
    for batch in range(round(n / batch_size)):
        start = batch * batch_size
        end = start + batch_size
        # perform update on a batch
        x_batch = x[start:end]
        y_batch = y[start:end]

        # build computation graph
        y_hat = a * x_batch + b
        myloss = torch.mean((y_batch - y_hat)**2)

        # gradient calculation
        myloss.backward()

        # SGD update
        with torch.no_grad():
            a -= lr * a.grad
            b -= lr * b.grad

        # avoid gradient accumulation
        a.grad.zero_()
        b.grad.zero_()

print(f"Epoch: {epoch + 1} / {nepochs}")
print(f"MSE: {myloss: .2e}; a: {a.item():.4f}; b: {b.item():.4f}")
```

1. Use a batch of data to
build a computation graph

Linear regression with SGD – The naive way

Output:

```
Epoch: 1 / 10
MSE: 3.69e-02; a: 1.4460; b: 1.2822
Epoch: 2 / 10
MSE: 9.24e-03; a: 1.7226; b: 1.1413
Epoch: 3 / 10
MSE: 2.32e-03; a: 1.8611; b: 1.0707
Epoch: 4 / 10
MSE: 5.81e-04; a: 1.9305; b: 1.0354
Epoch: 5 / 10
MSE: 1.46e-04; a: 1.9652; b: 1.0177
Epoch: 6 / 10
MSE: 3.65e-05; a: 1.9826; b: 1.0089
Epoch: 7 / 10
MSE: 9.15e-06; a: 1.9913; b: 1.0044
Epoch: 8 / 10
MSE: 2.29e-06; a: 1.9956; b: 1.0022
Epoch: 9 / 10
MSE: 5.75e-07; a: 1.9978; b: 1.0011
Epoch: 10 / 10
MSE: 1.44e-07; a: 1.9989; b: 1.0006
```





```
for epoch in range(nepochs):
    for batch in range(round(n / batch_size)):
        start = batch * batch_size
        end = start + batch_size
        # perform update on a batch
        x_batch = x[start:end]
        y_batch = y[start:end]

        # build computation graph
        y_hat = a * x_batch + b
        myloss = torch.mean((y_batch - y_hat)**2)

        # gradient calculation
        myloss.backward()

        # SGD update
        with torch.no_grad():
            a -= lr * a.grad
            b -= lr * b.grad

        # avoid gradient accumulation
        a.grad.zero_()
        b.grad.zero_()

print(f"Epoch: {epoch + 1} / {nepochs}")
print(f"MSE: {myloss: .2e}; a: {a.item():.4f}; b: {b.item():.4f}")
```

2. Calculate gradient based
on the computation graph

Linear regression with SGD – The naive way

Output:

```
Epoch: 1 / 10
MSE: 3.69e-02; a: 1.4460; b: 1.2822
Epoch: 2 / 10
MSE: 9.24e-03; a: 1.7226; b: 1.1413
Epoch: 3 / 10
MSE: 2.32e-03; a: 1.8611; b: 1.0707
Epoch: 4 / 10
MSE: 5.81e-04; a: 1.9305; b: 1.0354
Epoch: 5 / 10
MSE: 1.46e-04; a: 1.9652; b: 1.0177
Epoch: 6 / 10
MSE: 3.65e-05; a: 1.9826; b: 1.0089
Epoch: 7 / 10
MSE: 9.15e-06; a: 1.9913; b: 1.0044
Epoch: 8 / 10
MSE: 2.29e-06; a: 1.9956; b: 1.0022
Epoch: 9 / 10
MSE: 5.75e-07; a: 1.9978; b: 1.0011
Epoch: 10 / 10
MSE: 1.44e-07; a: 1.9989; b: 1.0006
```





```
for epoch in range(nepochs):
    for batch in range(round(n / batch_size)):
        start = batch * batch_size
        end = start + batch_size
        # perform update on a batch
        x_batch = x[start:end]
        y_batch = y[start:end]

        # build computation graph
        y_hat = a * x_batch + b
        myloss = torch.mean((y_batch - y_hat)**2)

        # gradient calculation
        myloss.backward()

        # SGD update
        with torch.no_grad():
            a -= lr * a.grad
            b -= lr * b.grad

        # avoid gradient accumulation
        a.grad.zero_()
        b.grad.zero_()

print(f"Epoch: {epoch + 1} / {nepochs}")
print(f"MSE: {myloss: .2e}; a: {a.item():.4f}; b: {b.item():.4f}")
```

3. Update parameters
and zero out gradients

Linear regression with SGD – The naive way

Output:

```
Epoch: 1 / 10
MSE: 3.69e-02; a: 1.4460; b: 1.2822
Epoch: 2 / 10
MSE: 9.24e-03; a: 1.7226; b: 1.1413
Epoch: 3 / 10
MSE: 2.32e-03; a: 1.8611; b: 1.0707
Epoch: 4 / 10
MSE: 5.81e-04; a: 1.9305; b: 1.0354
Epoch: 5 / 10
MSE: 1.46e-04; a: 1.9652; b: 1.0177
Epoch: 6 / 10
MSE: 3.65e-05; a: 1.9826; b: 1.0089
Epoch: 7 / 10
MSE: 9.15e-06; a: 1.9913; b: 1.0044
Epoch: 8 / 10
MSE: 2.29e-06; a: 1.9956; b: 1.0022
Epoch: 9 / 10
MSE: 5.75e-07; a: 1.9978; b: 1.0011
Epoch: 10 / 10
MSE: 1.44e-07; a: 1.9989; b: 1.0006
```



Linear regression with SGD

– The PyTorch way

Built-in functionalities

- Dataset preparation:
 - [torch.utils.data](#)
 - Check [Datasets & DataLoaders](#)
- Define training model:
 - Check [torch.nn](#) for all kinds of components to build your own model
- Optimization algorithms:
 - Check [torch.optim](#) for various opt methods

```
from torch import nn, optim
from torch.utils.data import Dataset, DataLoader
```

```
class SimpleLinear(Dataset):
    def __init__(self, n, a, b):
        self.n = n
        self.x = torch.rand(n, 1)
        self.y = a * self.x + b

    def __len__(self):
        return self.n

    def __getitem__(self, idx):
        return self.x[idx], self.y[idx]

# create a training model
# (cont.) also initializes a and b
mymodel = nn.Sequential(nn.Linear(1, 1))

# prepare dataset and dataloaders
mydata = SimpleLinear(n = 2000, a = 2, b = 1)
mydataloader = DataLoader(mydata, batch_size = 200)

# set up optimization with SGD
criterion = nn.MSELoss()
optimizer = optim.SGD(mymodel.parameters(), lr = 0.5)
```



Linear regression with SGD

– The PyTorch way

Prepare our datasets

- Dataset (a Python class)
 - `__init__`: initialize the dataset
 - `__len__`: sample size of the dataset
 - `__getitem__`: fetch a sample with `idx`
 - As the input of `DataLoader` function
- `DataLoader` (a Python function)
 - Divide the dataset into batches
 - Check [this](#) for more advanced usages

```
class SimpleLinear(Dataset):
    def __init__(self, n, a, b):
        self.n = n
        self.x = torch.rand(n, 1)
        self.y = a * self.x + b

    def __len__(self):
        return self.n

    def __getitem__(self, idx):
        return self.x[idx], self.y[idx]

# create a training model
# (cont.) also initializes a and b
mymodel = nn.Sequential(nn.Linear(1, 1))

# prepare dataset and dataloaders
mydata = SimpleLinear(n = 2000, a = 2, b = 1)
mydataloader = DataLoader(mydata, batch_size = 200)

# set up optimization with SGD
criterion = nn.MSELoss()
optimizer = optim.SGD(mymodel.parameters(), lr = 0.5)
```



Linear regression with SGD

– The PyTorch way

Create the training model:

- `nn.Sequential`
 - Build a model with sequential operations
 - `nn.Linear(m, n)`
 - A **linear operator** of shape $n \times m$
 - A **bias vector** of shape $n \times 1$ (default)
 - All parameters are **initialized automatically upon creation**
 - Check [torch.nn](https://pytorch.org/docs/nn.html) for other operations



```
class SimpleLinear(Dataset):
    def __init__(self, n, a, b):
        self.n = n
        self.x = torch.rand(n, 1)
        self.y = a * self.x + b

    def __len__(self):
        return self.n

    def __getitem__(self, idx):
        return self.x[idx], self.y[idx]

# create a training model
# (cont.) also initializes a and b
mymodel = nn.Sequential(nn.Linear(1, 1))

# prepare dataset and dataloaders
mydata = SimpleLinear(n = 2000, a = 2, b = 1)
mydataloader = DataLoader(mydata, batch_size = 200)

# set up optimization with SGD
criterion = nn.MSELoss()
optimizer = optim.SGD(mymodel.parameters(), lr = 0.5)
```



Linear regression with SGD

– The PyTorch way

Set up optimization:

- The loss function
- The optimization method
- Optimization parameters

```
class SimpleLinear(Dataset):  
    def __init__(self, n, a, b):  
        self.n = n  
        self.x = torch.rand(n, 1)  
        self.y = a * self.x + b  
  
    def __len__(self):  
        return self.n  
  
    def __getitem__(self, idx):  
        return self.x[idx], self.y[idx]  
  
# create a training model  
# (cont.) also initializes a and b  
mymodel = nn.Sequential(nn.Linear(1, 1))  
  
# prepare dataset and dataloaders  
mydata = SimpleLinear(n = 2000, a = 2, b = 1)  
mydataloader = DataLoader(mydata, batch_size = 200)  
  
# set up optimization with SGD  
criterion = nn.MSELoss()  
optimizer = optim.SGD(mymodel.parameters(), lr = 0.5)
```



Linear regression with SGD

– The PyTorch way



```
nepochs = 10
for epoch in range(nepochs):
    for x_batch, y_batch in mydataloader:
        # build computation graph
        yhat = mymodel(x_batch)
        myloss = criterion(y_batch, yhat)

        # optimization
        optimizer.zero_grad() # zero out gradient
        myloss.backward() # back propagation
        optimizer.step() # update parameters
```

Output:

```
Epoch: 1 / 10
MSE: 3.52e-02; a: 1.3904; b: 1.3368
Epoch: 2 / 10
MSE: 9.21e-03; a: 1.6882; b: 1.1722
Epoch: 3 / 10
MSE: 2.41e-03; a: 1.8406; b: 1.0881
Epoch: 4 / 10
MSE: 6.30e-04; a: 1.9185; b: 1.0450
Epoch: 5 / 10
MSE: 1.65e-04; a: 1.9583; b: 1.0230
Epoch: 6 / 10
MSE: 4.31e-05; a: 1.9787; b: 1.0118
Epoch: 7 / 10
MSE: 1.13e-05; a: 1.9891; b: 1.0060
Epoch: 8 / 10
MSE: 2.95e-06; a: 1.9944; b: 1.0031
Epoch: 9 / 10
MSE: 7.71e-07; a: 1.9971; b: 1.0016
Epoch: 10 / 10
MSE: 2.02e-07; a: 1.9985; b: 1.0008
```



Linear regression with SGD

– The PyTorch way



```
nepochs = 10
for epoch in range(nepochs):
    for x_batch, y_batch in mydataloader:
        # build computation graph
        yhat = mymodel(x_batch)
        myloss = criterion(y_batch, yhat)

        # optimization
        optimizer.zero_grad() # zero out gradient
        myloss.backward() # back propagation
        optimizer.step() # update parameters
```

The whole
optimization part

Output:

```
Epoch: 1 / 10
MSE: 3.52e-02; a: 1.3904; b: 1.3368
Epoch: 2 / 10
MSE: 9.21e-03; a: 1.6882; b: 1.1722
Epoch: 3 / 10
MSE: 2.41e-03; a: 1.8406; b: 1.0881
Epoch: 4 / 10
MSE: 6.30e-04; a: 1.9185; b: 1.0450
Epoch: 5 / 10
MSE: 1.65e-04; a: 1.9583; b: 1.0230
Epoch: 6 / 10
MSE: 4.31e-05; a: 1.9787; b: 1.0118
Epoch: 7 / 10
MSE: 1.13e-05; a: 1.9891; b: 1.0060
Epoch: 8 / 10
MSE: 2.95e-06; a: 1.9944; b: 1.0031
Epoch: 9 / 10
MSE: 7.71e-07; a: 1.9971; b: 1.0016
Epoch: 10 / 10
MSE: 2.02e-07; a: 1.9985; b: 1.0008
```



Linear regression with SGD

– Comparison of two ways

The PyTorch way

```
▶ epochs = 10
for epoch in range(epochs):
    for x_batch, y_batch in mydataloader:
        # build computation graph
        yhat = mymodel(x_batch)
        myloss = criterion(y_batch, yhat)

        # optimization
        optimizer.zero_grad() # zero out gradient
        myloss.backward() # back propagation
        optimizer.step() # update parameters
```



Agnostic to datasets, models,
losses and optimization methods!

The naive way

```
▶ for epoch in range(nepochs):
    for batch in range(round(n / batch_size)):
        start = batch * batch_size
        end = start + batch_size
        # perform update on a batch
        x_batch = x[start:end]
        y_batch = y[start:end]

        # build computation graph
        y_hat = a * x_batch + b
        myloss = torch.mean((y_batch - y_hat)**2)

        # gradient calculation
        myloss.backward()

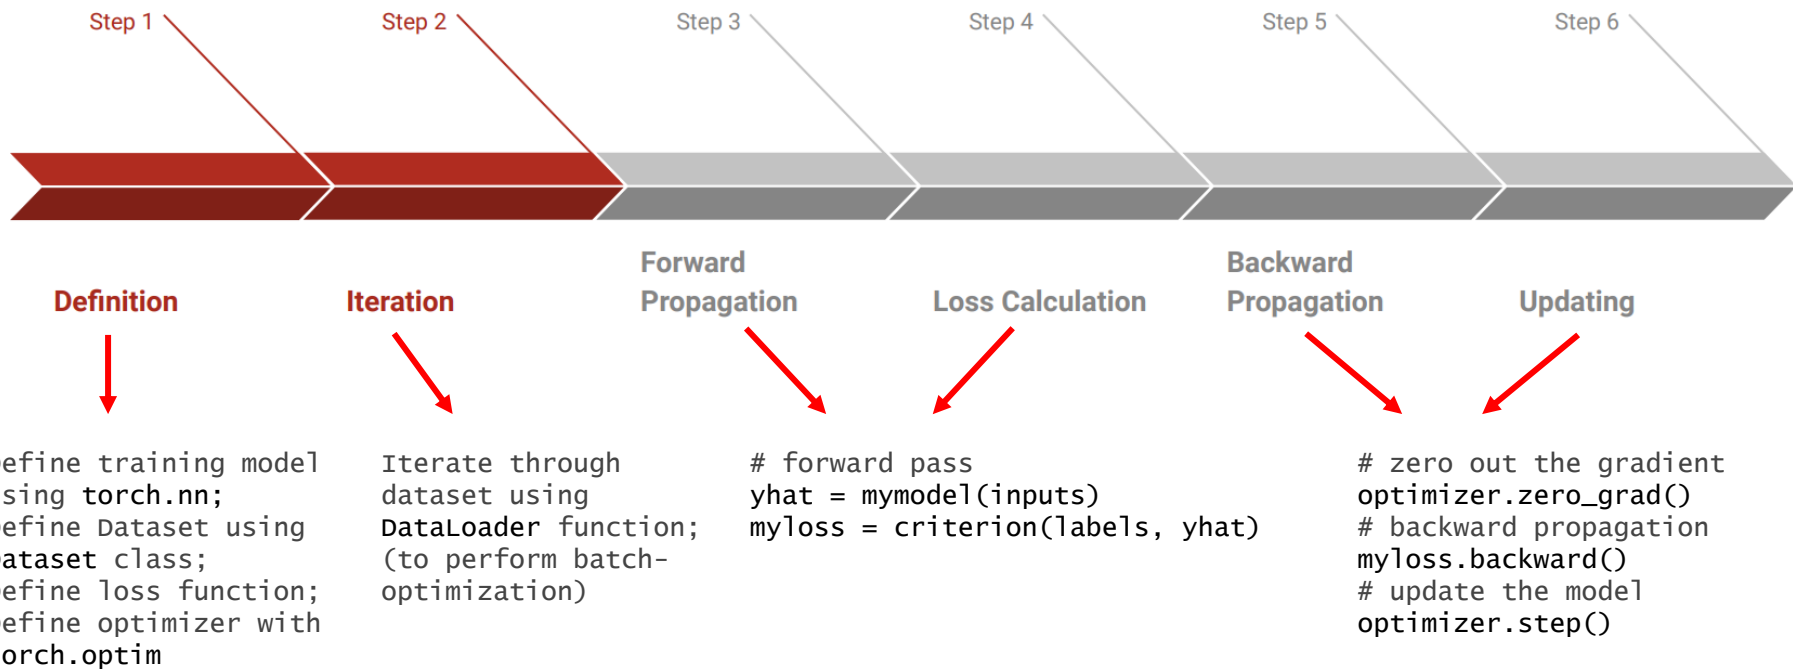
        # SGD update
        with torch.no_grad():
            a -= lr * a.grad
            b -= lr * b.grad

        # avoid gradient accumulation
        a.grad.zero_()
        b.grad.zero_()
```



ML workflow in PyTorch

Credit: HPRC Short Course by Jian Tao, TAMU



Outline

- Set up working environment
- Introduction to PyTorch framework
- Deep learning with PyTorch
 - Regression using shallow ReLU net
 - Image classification with CNN
- Additional resources



From linear model to shallow ReLU network

Simple linear model

```
# linear model
LinearModel = nn.Sequential(nn.Linear(1, 1))
print(f"Slope: {mymodel[0].weight.item(): .4f}")
print(f"Intercept: {mymodel[0].bias.item(): .4f}")
```

```
➡ Slope: 0.0941
Intercept: -0.2586
```

Shallow ReLU neural net

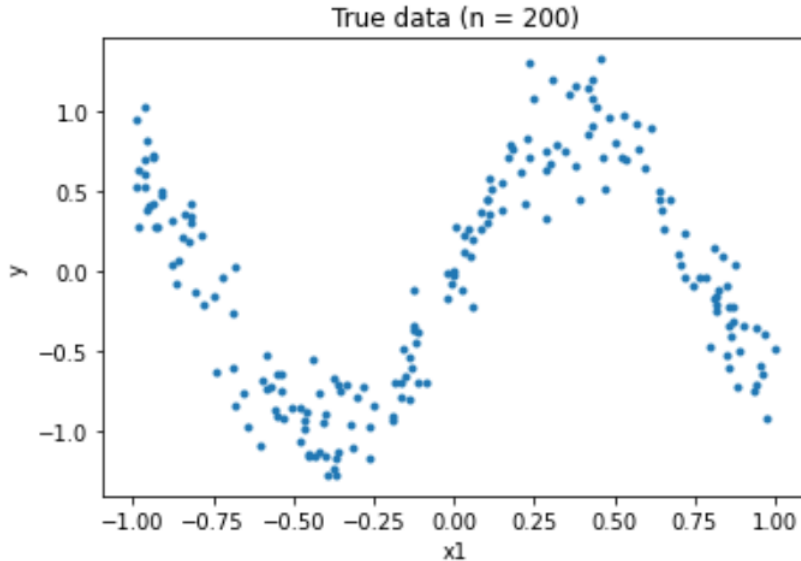
```
# shallow ReLU net
ReLUnn2 = nn.Sequential(
    nn.Linear(1, 3),
    nn.ReLU(),
    nn.Linear(3, 1)
)
for param in ReLUnn2.parameters():
    print(type(param.data), param.size())
```

```
➡ <class 'torch.Tensor'> torch.Size([3, 1])
<class 'torch.Tensor'> torch.Size([3])
<class 'torch.Tensor'> torch.Size([1, 3])
<class 'torch.Tensor'> torch.Size([1])
```



Fit a shallow ReLU neural network to Sine function

What dose the data look like?



- x is 5-dimensional, but only the first coordinate x_1 is relevant to y (sparse)

What is our training model?



```
# initialize a training model
```

```
train_model = nn.Sequential(  
    nn.Linear(5, 256),  
    nn.ReLU(),  
    nn.Linear(256, 1)  
)  
train_model
```



```
Sequential(  
  (0): Linear(in_features=5, out_features=256, bias=True)  
  (1): ReLU()  
  (2): Linear(in_features=256, out_features=1, bias=True)  
)
```

- A shallow ReLU neural network with 256 hidden nodes



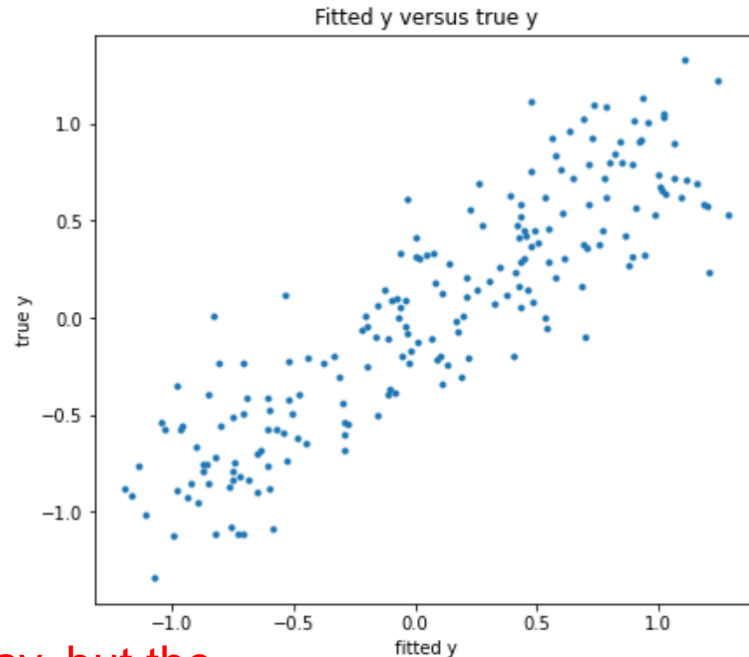

Fit a shallow ReLU neural network to Sine function

Trained model on **TRAINING** data



Fit a shallow ReLU neural network to Sine function

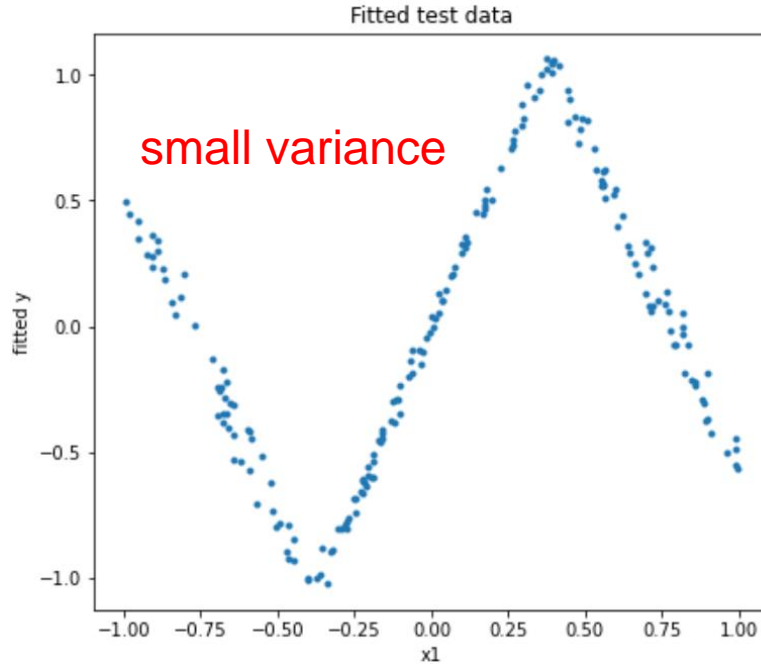
Trained model on **TESTING** data




Seems okay, but the
Variance is a little bit large



Fit a shallow ReLU neural network to Sine function



Trained model on **TESTING** data and with **REGULARIZATION**



Code: `optimizer = optim.Adam(train_model.parameters(), lr = 0.001, weight_decay = 0.01)`



Regularization of neural net

- To achieve better generalization

- Explicit regularization:


$$\min_{\mathbf{W}} \frac{1}{m} \sum_{i=1}^m \ell(\mathbf{y}_i, \text{DNN}_{\mathbf{W}}(\mathbf{x}_i)) + \lambda \Omega(\mathbf{W})$$

- Norms of weight matrices
- Norms of gradients / Jacobians
- ...

- Miscellaneous: A neural net structure can also induce regularization. E.g. The convolutional neural net (CNN)

- Implicit regularization:

- The regularization that is not built in the objective function (the loss)
- Regularization induced by an **optimization algorithm**: SGD tends to find a solution with small norm (regularized solution)
- **Early stopping**; batch normalization; dropout



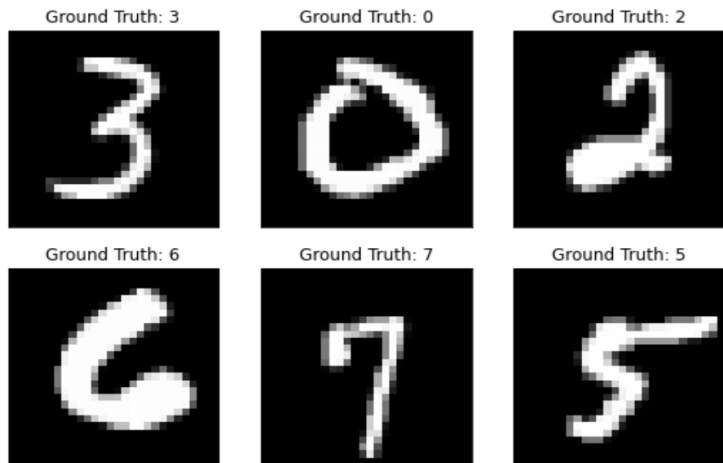
Credit: CSCI 5980/8980: Think Deep Learning offered by Ju Sun, UMN



Image classification using CNN

Dataset: MNIST

- One of the benchmark datasets of deep learning for image classification
- Classifying digits as 0, 1, ..., 9



Our CNN model

```
def forward(self, x):  
    x = self.conv1(x)  
    x = F.relu(x)  
    x = self.conv2(x)  
    x = F.relu(x)  
    x = F.max_pool2d(x, 2)  
    x = self.dropout1(x)  
    x = torch.flatten(x, 1)  
    x = self.fc1(x)  
    x = F.relu(x)  
    x = self.dropout2(x)  
    x = self.fc2(x)  
    output = F.log_softmax(x, dim=1)  
    return output
```

Image classification using CNN

Our CNN model

Two convolutional layers

```
def forward(self, x):  
    x = self.conv1(x)  
    x = F.relu(x)  
    x = self.conv2(x)  
    x = F.relu(x)  
    x = F.max_pool2d(x, 2)  
    x = self.dropout1(x)  
    x = torch.flatten(x, 1)  
    x = self.fc1(x)  
    x = F.relu(x)  
    x = self.dropout2(x)  
    x = self.fc2(x)  
    output = F.log_softmax(x, dim=1)  
    return output
```



Image classification using CNN

Our CNN model

One max-pooling layer

```
def forward(self, x):  
    x = self.conv1(x)  
    x = F.relu(x)  
    x = self.conv2(x)  
    x = F.relu(x)  
    x = F.max_pool2d(x, 2)  
    x = self.dropout1(x)  
    x = torch.flatten(x, 1)  
    x = self.fc1(x)  
    x = F.relu(x)  
    x = self.dropout2(x)  
    x = self.fc2(x)  
    output = F.log_softmax(x, dim=1)  
    return output
```



Image classification using CNN

Our CNN model

```
def forward(self, x):  
    x = self.conv1(x)  
    x = F.relu(x)  
    x = self.conv2(x)  
    x = F.relu(x)  
    x = F.max_pool2d(x, 2)  
    x = self.dropout1(x)  
    x = torch.flatten(x, 1)  
    x = self.fc1(x)  
    x = F.relu(x)  
    x = self.dropout2(x)  
    x = self.fc2(x)  
    output = F.log_softmax(x, dim=1)  
    return output
```

Two fully connected layers




Image classification using CNN

Our CNN model

ReLU activations

```
def forward(self, x):  
    x = self.conv1(x)  
    x = F.relu(x)  
    x = self.conv2(x)  
    x = F.relu(x)  
    x = F.max_pool2d(x, 2)  
    x = self.dropout1(x)  
    x = torch.flatten(x, 1)  
    x = self.fc1(x)  
    x = F.relu(x)  
    x = self.dropout2(x)  
    x = self.fc2(x)  
    output = F.log_softmax(x, dim=1)  
    return output
```

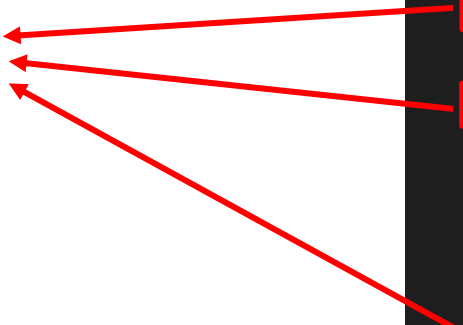




Image classification using CNN

Dropout for regularization:

- randomly kills inner neurons with some probability p



(a) Standard Neural Net



(b) After applying dropout.

Credit: Srivastava, N., Hinton, G., Krizhevsky, A., Sutskever, I., & Salakhutdinov, R. (2014). Dropout: a simple way to prevent neural networks from overfitting. *The journal of machine learning research*, 15(1), 1929-1958.

Our CNN model

```
def forward(self, x):  
    x = self.conv1(x)  
    x = F.relu(x)  
    x = self.conv2(x)  
    x = F.relu(x)  
    x = F.max_pool2d(x, 2)  
    x = self.dropout1(x)  
    x = torch.flatten(x, 1)  
    x = self.fc1(x)  
    x = F.relu(x)  
    x = self.dropout2(x)  
    x = self.fc2(x)  
    output = F.log_softmax(x, dim=1)  
    return output
```



Image classification using CNN

The total trainable
parameters is 1,199,882!

Layer (type)	Input Shape	Param #	Tr. Param #
Conv2d-1	[1, 1, 28, 28]	320	320
Conv2d-2	[1, 32, 26, 26]	18,496	18,496
Dropout-3	[1, 64, 12, 12]	0	0
Linear-4	[1, 9216]	1,179,776	1,179,776
Dropout-5	[1, 128]	0	0
Linear-6	[1, 128]	1,290	1,290
Total params: 1,199,882			
Trainable params: 1,199,882			
Non-trainable params: 0			



Check GPU availability

CPU



```
# check device availability  
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")  
print(device)
```

cpu

GPU



```
# check device availability  
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")  
print(device)
```



cuda



Download and preprocess the MNIST training / testing dataset

Other transforms:
cropping, translation,
rotation, padding...

```
# define image transformation
transform=transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize((0.1307,), (0.3081,))
])

# import MNIST dataset
dataset1 = datasets.MNIST('../data', train = True, download = True, transform = transform)
dataset2 = datasets.MNIST('../data', train = False, transform = transform)

# only use 2000 training images and 20 testing images
train_index = np.random.choice(range(60000), 2000, replace = False)
test_index = np.random.choice(range(10000), 20, replace = False)
train_data = Subset(dataset1, train_index)
test_data = Subset(dataset2, test_index)
```



Download and preprocess the MNIST training / testing dataset

Most popular benchmark datasets can be loaded via [torchvision](#) library

```
# define image transformation
transform=transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize((0.1307,), (0.3081,))
])

# import MNIST dataset
dataset1 = datasets.MNIST('../data', train = True, download = True, transform = transform)
dataset2 = datasets.MNIST('../data', train = False, transform = transform)

# only use 2000 training images and 20 testing images
train_index = np.random.choice(range(60000), 2000, replace = False)
test_index = np.random.choice(range(10000), 20, replace = False)
train_data = Subset(dataset1, train_index)
test_data = Subset(dataset2, test_index)
```

and don't forget to prepare DataLoader for training



Define a training utils function

```
# a utils function for training
def train(model, device, train_loader, optimizer, epoch):
    model.train() # enable dropout
    correct = 0
    for batch_idx, (data, target) in enumerate(train_loader):
        # transfer batches of data to specified device
        data, target = data.to(device), target.to(device)

        # build computation graph
        output = model(data)
        loss = F.nll_loss(output, target)

        # the optimization part
        optimizer.zero_grad()
        loss.backward()
        optimizer.step()

    # count the corrected classified cases
    pred = output.argmax(dim = 1, keepdim = True) # get the index of the max log-probability
    correct += pred.eq(target.view_as(pred)).sum().item()

    print(f"Epoch {epoch + 1}:")
    print(f"Training accuracy: {100. * correct / len(train_loader.dataset): .2f} %")
```

Similar as before,
nothing new



Train our CNN



```
model = Net().to(device)
optimizer = optim.Adam(model.parameters(), lr = 0.005)

nepochs = 10
for epoch in range(nepochs):
    train(model, device, train_loader, optimizer, epoch)
    test(model, device, test_loader)
```



```
Epoch 1: Training accuracy: 67.10 %; Testing accuracy: 19/20 (95%)
Epoch 2: Training accuracy: 88.30 %; Testing accuracy: 20/20 (100%)
Epoch 3: Training accuracy: 92.85 %; Testing accuracy: 20/20 (100%)
Epoch 4: Training accuracy: 94.75 %; Testing accuracy: 20/20 (100%)
Epoch 5: Training accuracy: 94.90 %; Testing accuracy: 20/20 (100%)
Epoch 6: Training accuracy: 95.95 %; Testing accuracy: 20/20 (100%)
Epoch 7: Training accuracy: 96.30 %; Testing accuracy: 20/20 (100%)
Epoch 8: Training accuracy: 96.70 %; Testing accuracy: 20/20 (100%)
Epoch 9: Training accuracy: 97.20 %; Testing accuracy: 19/20 (95%)
Epoch 10: Training accuracy: 97.10 %; Testing accuracy: 19/20 (95%)
```



Save / load the trained model

- State dictionary method: only save parameter values:



```
# save parameter values
torch.save(model.state_dict(), "model_dict.pt")

# load parameter values
model1 = Net()
model1.load_state_dict(torch.load("model_dict.pt"))
model1.eval()
```



```
Net(
  (conv1): Conv2d(1, 32, kernel_size=(3, 3), stride=(1, 1))
  (conv2): Conv2d(32, 64, kernel_size=(3, 3), stride=(1, 1))
  (dropout1): Dropout(p=0.25, inplace=False)
  (dropout2): Dropout(p=0.5, inplace=False)
  (fc1): Linear(in_features=9216, out_features=128, bias=True)
  (fc2): Linear(in_features=128, out_features=10, bias=True)
)
```



Save / load the trained model

- Checkpoint method: save for resuming training later

```
# save checkpoint for resuming training
torch.save({
    'epoch': epoch,
    'model_state_dict': model.state_dict(),
    'optimizer_state_dict': optimizer.state_dict()
}, "model_checkpoint.pt")

# initialize class and optimizer
model_new = Net()
optimizer_new = optim.Adam(model_new.parameters(), lr = 0.005)

# load checkpoint
checkpoint = torch.load("model_checkpoint.pt")
model_new.load_state_dict(checkpoint['model_state_dict'])
optimizer_new.load_state_dict(checkpoint['optimizer_state_dict'])
epoch = checkpoint['epoch']

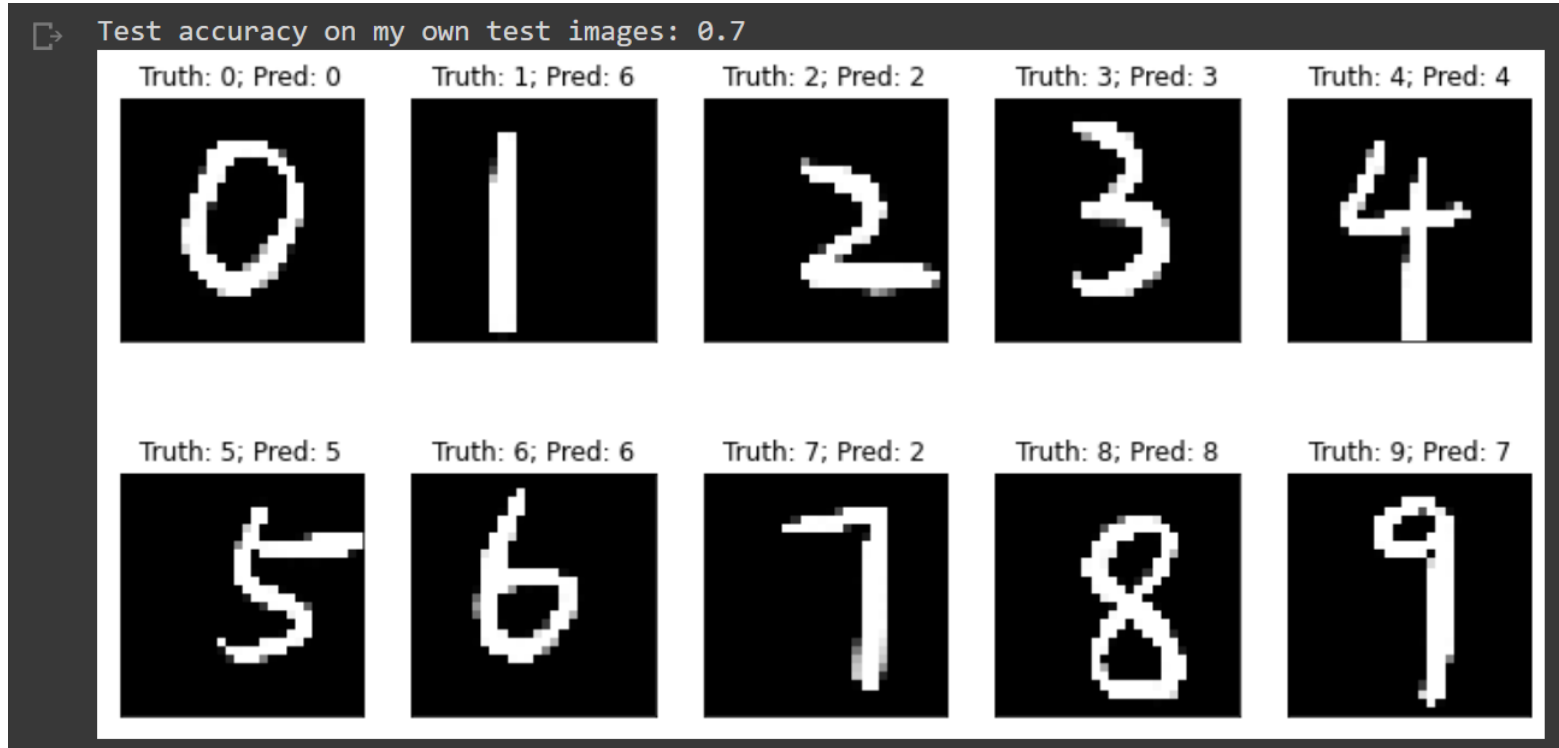
# do whatever you like :)
model_new.eval()
# - or -
model_new.train()
```

Check [saving and loading models](#) for more details



Robustness issue of CNN

Solution: data augmentation!



Robustness issue of CNN

Solution: adversarial training!




x

“panda”

57.7% confidence

$+ .007 \times$



$\text{sign}(\nabla_x J(\theta, x, y))$

“nematode”

8.2% confidence

$=$



$x + \epsilon \text{sign}(\nabla_x J(\theta, x, y))$

“gibbon”

99.3 % confidence

Credit: Goodfellow, I. J., Shlens, J., & Szegedy, C. (2014). Explaining and harnessing adversarial examples. arXiv preprint arXiv:1412.6572.



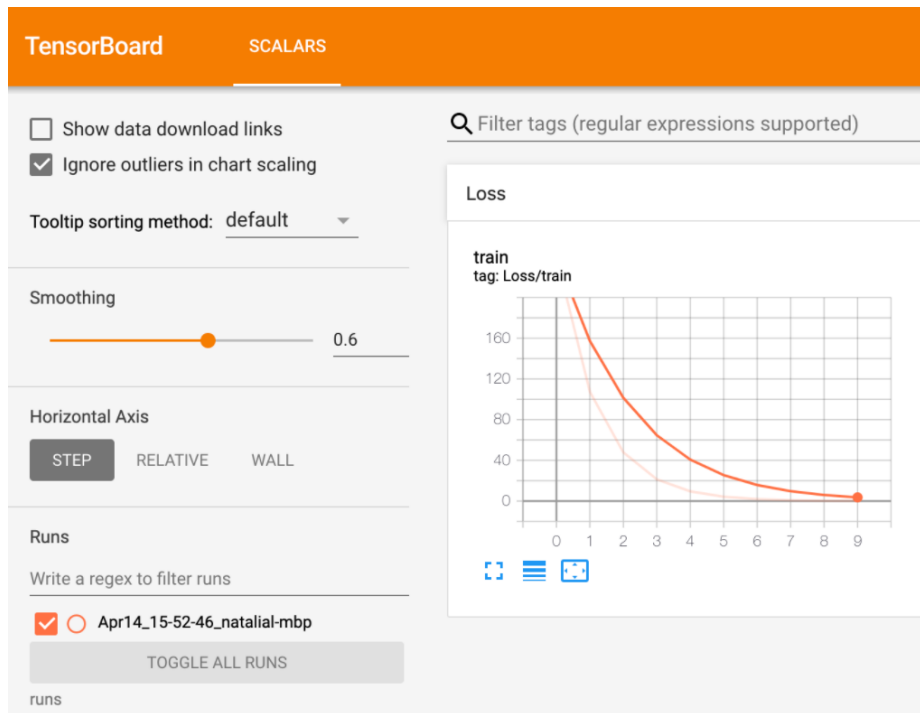
Outline

- Set up working environment
- Introduction to PyTorch framework
- Deep learning with PyTorch
- Other resources



Visualization with TensorBoard

- Track and visualize:
 - Metrics such as loss and accuracy
 - The model graph
 - Histograms, images and much more
- Some resources:
 - [How to use TensorBoard with PyTorch](#)
 - [Visualizing models, data, and training with TensorBoard](#)
 - [PyTorch TensorBoard support](#)



Tutorials and courses:

Books and tutorials:

- [Dive into Deep Learning](#) (livebook)
- [Deep Learning](#) by Ian Goodfellow
- [Official PyTorch tutorial](#)
- [Deep learning with Python](#) (livebook)
- [UvA DL Notebooks](#)

Courses:

- [DL/ML tutorial](#) by Hung-Yi Lee
- [Deep learning](#) course by Yann LeCun
- [Deep learning with Pytorch](#)
- [Stanford STAT385 series](#)
- [Think Deep Learning](#) by Ju Sun





UNIVERSITY OF MINNESOTA

Driven to Discover®

Crookston Duluth Morris Rochester Twin Cities

The University of Minnesota is an equal opportunity educator and employer.